



# HTTP/2

---

Frederik Deweerdt - @fdeweerdt



**89.7%**

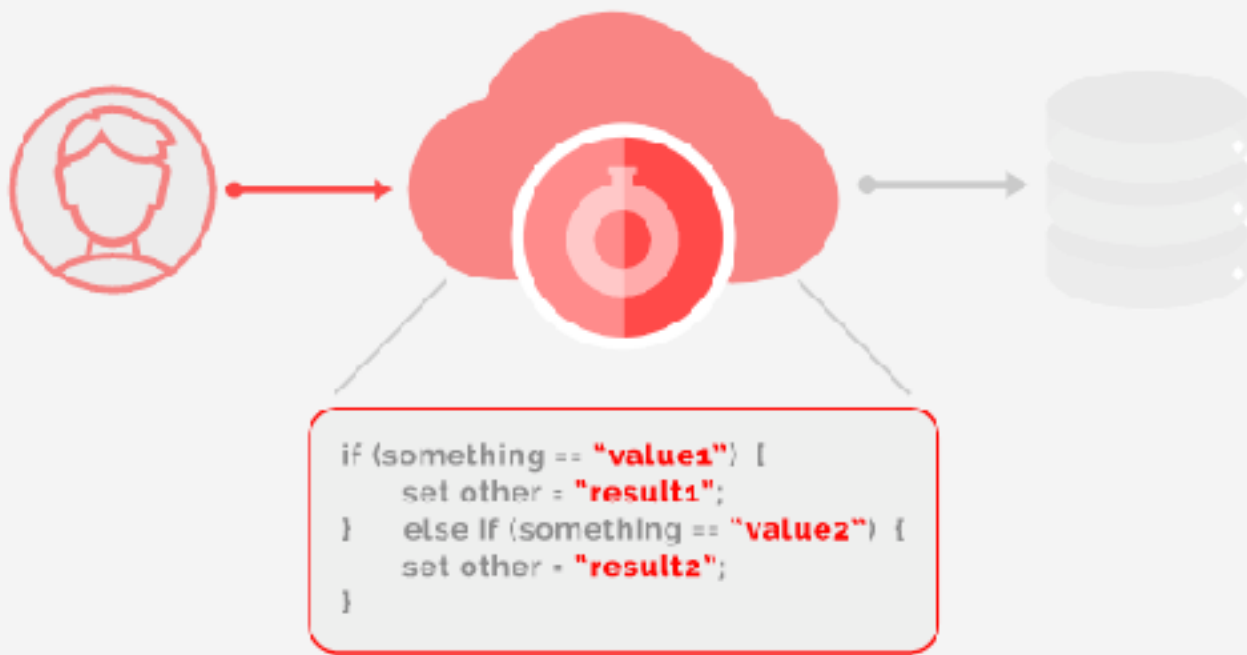
Cache hit ratio

**0.090 ms**

Time to 1st byte

**2.78 Tb/s**

Bandwidth



```
sub vcl_recv {  
    #FASTLY recv  
  
    if (fastly_info.is_h2 && req.url ~ "^/index.html")  
    {  
        h2.push("css1.css");  
    }  
}
```

# H2O

the optimized HTTP/1.x, HTTP/2 server

Star 4,801

Search

Top

Install

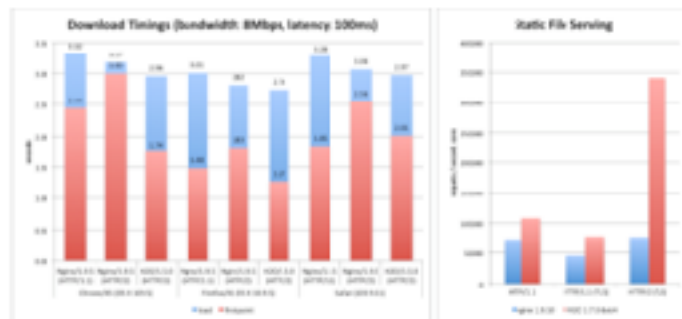
Configure

FAQ

Blog

Source

H2O is a new generation HTTP server that **provides quicker response to users with less CPU utilization** when compared to older generation of web servers. Designed from ground-up, the server takes full advantage of [HTTP/2](#) features including [prioritized content serving](#) and [server push](#), promising outstanding experience to the visitors of your web site.



Explanation of the benchmark charts can be found in the [benchmarks](#) page.

## Key Features

- HTTP/1.0, HTTP/1.1
- [HTTP/2](#)
  - full support for dependency and weight-based prioritization with [server-side tweaks](#)
  - [cache-aware server push](#)
- TCP Fast Open
- [TLS](#)
  - session resumption (standalone & memcached)
  - session tickets with automatic key rollover

# HTTP/2

# Why?

# RFC 2616 - 1999

---

Network Working Group  
Request for Comments: 2616  
Obsoletes: 2058  
Category: Standards Track

R. Fielding  
DC Irvine  
J. Gettys  
Compaq/WEC  
J. Mogul  
Compaq  
H. Frystyk  
NSC/NTT  
L. Masinter  
Xerox  
P. Leach  
Microsoft  
T. Berners-Lee  
NSC/NTT  
June 1999

## HyperText Transfer Protocol -- HTTP/1.1

### Status of this Memo

This document specifies an Internet standards track protocol for the Internet community, and requests discussion and suggestions for improvements. Please refer to the current edition of the "Internet Official Protocol Standards" (STD 1) for the standardization state and status of this protocol. Distribution of this memo is unlimited.

### Copyright Notice

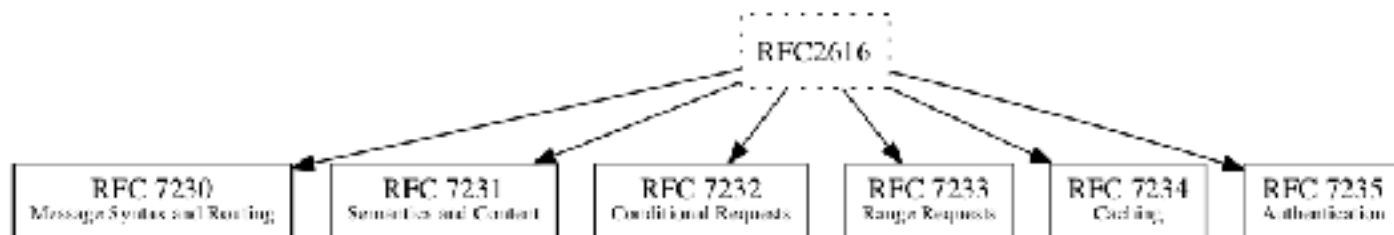
Copyright (C) The Internet Society (1999). All Rights Reserved.

### Abstract



# HTTP refresh - 2014

---



# HTTP and TCP

---

- HTTP/1.0: open - fetch - close
- HTTP/1.1: open - fetch - wait - fetch - wait  
- ... - close

# HTTP/1.1 concurrency

---

## RFC 2616 - Section 8.1.4

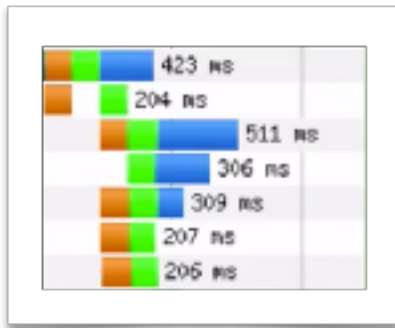
Clients that use persistent connections SHOULD limit the number of simultaneous connections that they maintain to a given server. A single-user client SHOULD NOT maintain more than 2 connections with any server or proxy. A proxy SHOULD use up to  $2*N$  connections to another server or proxy, where  $N$  is the number of simultaneously active users. These guidelines are intended to improve HTTP response times and avoid congestion.

# HTTP/1.1 concurrency

---

## RFC 2616 - Section 8.1.4

Clients that use persistent connections SHOULD limit the number of simultaneous connections that they maintain to a given server. A single-user client SHOULD NOT maintain more than 2 connections with any server or proxy. A proxy SHOULD use up to  $2*N$  connections to another server or proxy, where  $N$  is the number of simultaneously active users. These guidelines are intended to improve HTTP response times and avoid congestion.



# HTTP/1.1 concurrency

gopherfiles?latency=1000	200	http/1.1	document	Other	18.1 KB	404 ms	100 ms
gopherfiles?c=1&y=0&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	815 B	1.35 s	
gopherfiles?c=5&y=0&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	821 B	1.35 s	
gopherfiles?c=7&y=0&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	829 B	1.35 s	
gopherfiles?c=4&y=0&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	834 B	1.35 s	
gopherfiles?c=5&y=0&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	823 B	1.35 s	
gopherfiles?c=0&y=0&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	816 B	1.38 s	
gopherfiles?c=0&y=1&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	828 B	2.16 s	
gopherfiles?c=6&y=5&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	1.1 KB	2.17 s	
gopherfiles?c=7&y=2&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	1.0 KB	2.17 s	
gopherfiles?c=1&y=8&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	1.1 KB	2.18 s	
gopherfiles?c=10&y=2&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	826 B	2.18 s	
gopherfiles?c=10&y=5&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	829 B	2.18 s	
gopherfiles?c=8&y=2&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	829 B	3.24 s	
gopherfiles?c=13&y=5&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	830 B	3.24 s	
gopherfiles?c=4&y=1&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	1.0 KB	3.24 s	
gopherfiles?c=0&y=8&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	909 B	3.24 s	
gopherfiles?c=8&y=2&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	960 B	3.25 s	
gopherfiles?c=6&y=0&cachebuster=147153890..	200	http/1.1	image	gopherfiles?latency=10..	637 B	3.25 s	

# HTTP/1.1 concurrency

gopherfiles?latency=1000	200	http/1.1	document	Other	18.1 KB	404 ms	1.38 s
gopherfiles?c=1&y=0&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	815 B	1.35 s	
gopherfiles?c=5&y=0&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	821 B	1.35 s	
gopherfiles?c=7&y=0&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	829 B	1.35 s	
gopherfiles?c=4&y=0&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	834 B	1.35 s	
gopherfiles?c=5&y=0&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	823 B	1.35 s	
gopherfiles?c=0&y=0&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	816 B	1.38 s	
gopherfiles?c=0&y=1&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	828 B	2.16 s	
gopherfiles?c=6&y=5&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	1.1 KB	2.17 s	
gopherfiles?c=7&y=2&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	1.0 KB	2.17 s	
gopherfiles?c=1&y=8&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	1.1 KB	2.18 s	
gopherfiles?c=10&y=2&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	826 B	2.18 s	
gopherfiles?c=10&y=5&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	829 B	2.18 s	
gopherfiles?c=8&y=2&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	829 B	3.24 s	
gopherfiles?c=13&y=5&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	830 B	3.24 s	
gopherfiles?c=4&y=1&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	1.0 KB	3.24 s	
gopherfiles?c=0&y=8&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	909 B	3.24 s	
gopherfiles?c=8&y=2&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	960 B	3.25 s	
gopherfiles?c=6&y=0&cachebuster=147153890..	200	http/1.1	jpeg	gopherfiles?latency=10..	637 B	3.25 s	

# HTTP/1.1 pipelining

---

Fielding, et al.

Standards Track

[Page 45]

RFC 2616

HTTP/1.1

June 1999

## 8.1.2.2 Pipelining

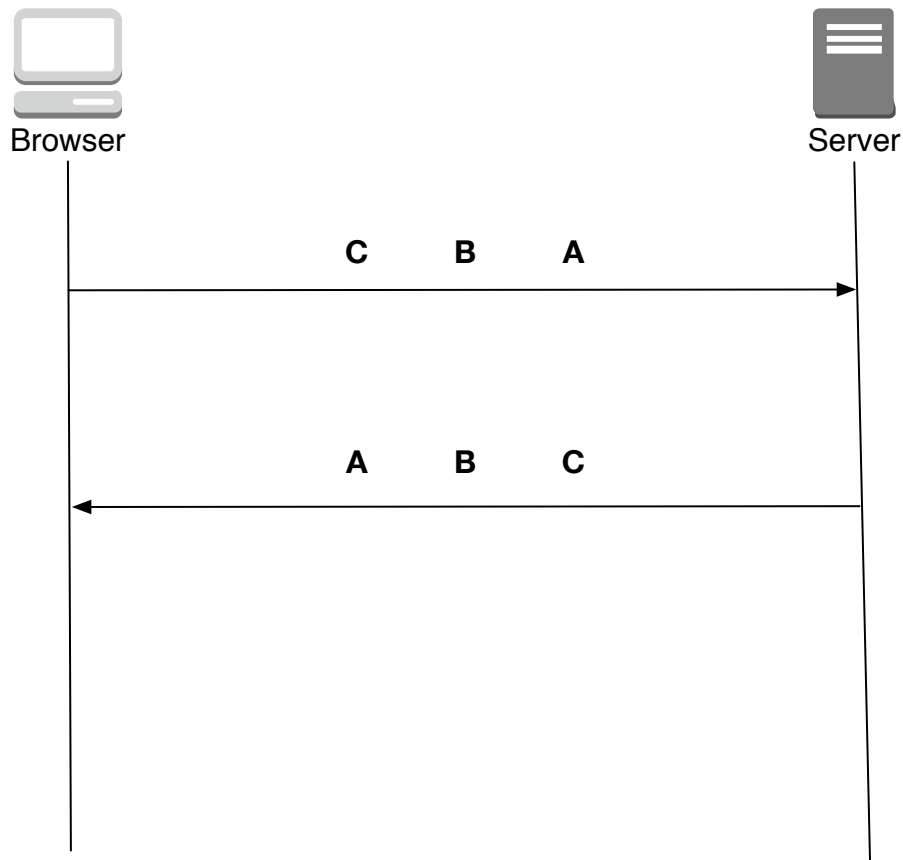
A client that supports persistent connections MAY "pipeline" its requests (i.e., send multiple requests without waiting for each response). A server MUST send its responses to those requests in the same order that the requests were received.

Clients which assume persistent connections and pipeline immediately after connection establishment SHOULD be prepared to retry their connection if the first pipelined attempt fails. If a client does such a retry, it MUST NOT pipeline before it knows the connection is persistent. Clients MUST also be prepared to resend their requests if the server closes the connection before sending all of the corresponding responses.

Clients SHOULD NOT pipeline requests using non-idempotent methods or non-idempotent sequences of methods (see section 9.1.2). Otherwise, a premature termination of the transport connection could lead to indeterminate results. A client wishing to send a non-idempotent request SHOULD wait to send that request until it has received the response status for the previous request.

# HTTP/1.1 pipelining

---

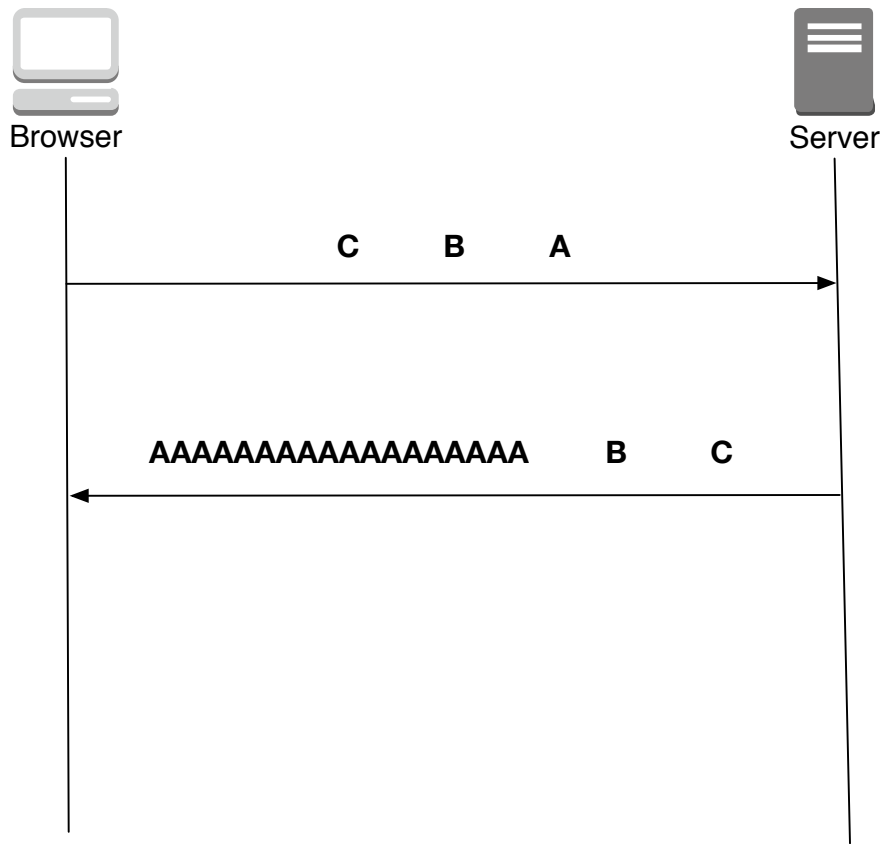




# Head of line blocking

# HOL blocking

---



# HOL blocking



Browser



Server

C B A



B C

# Pipelining improvements

---

Network Working Group  
Internet-Draft  
Expires: 1 November 2001

Jeffrey Mogul, Google AG  
6 April 2001

## Support for out-of-order responses in HTTP

draft-mogul-http-ooo-00.txt

### STATUS OF THIS MEMO

This document is an Internet-Draft and is in full conformance with all provisions of Section 10 of RFC2026.

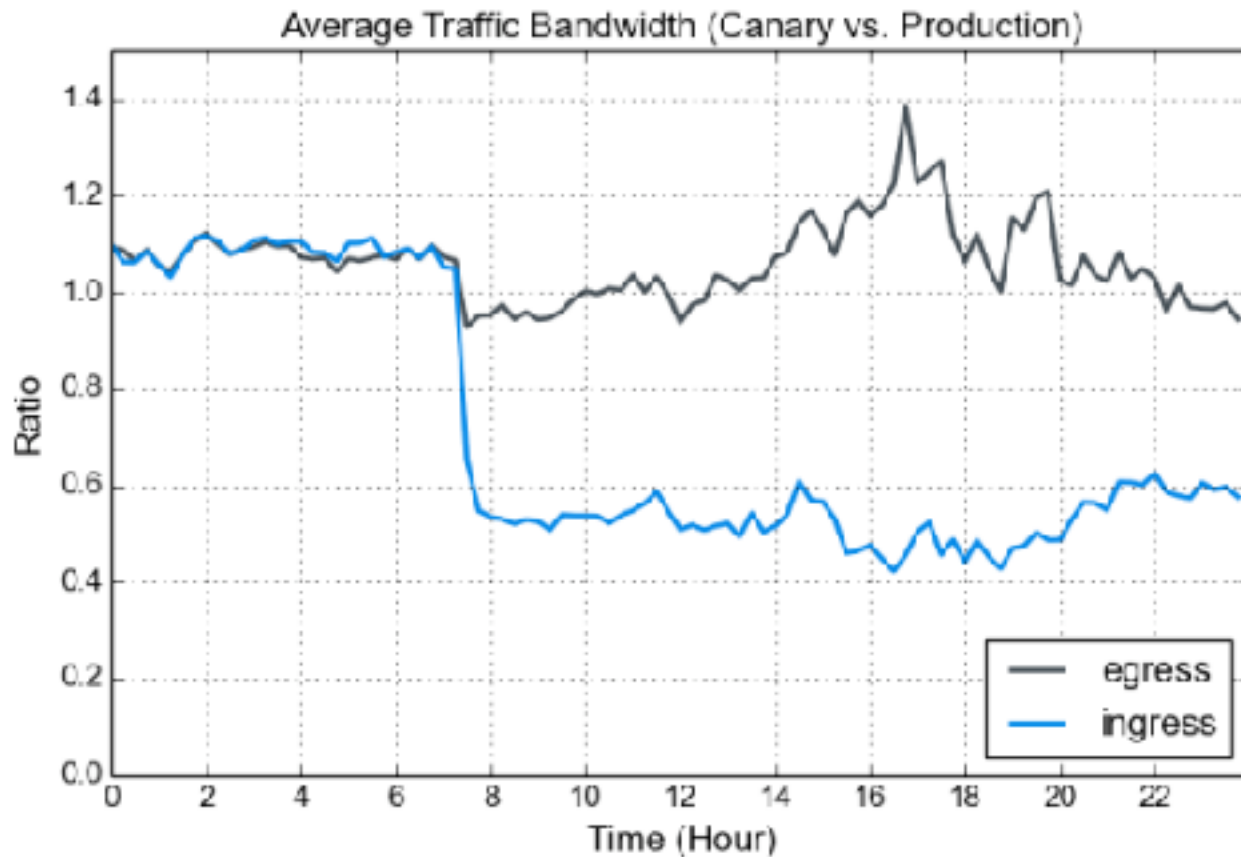
Internet-Drafts are working documents of the Internet Engineering Task Force (IETF), its areas, and its working groups. Note that other groups may also distribute working documents as Internet-Drafts.

Internet-Drafts are draft documents valid for a maximum of six months and may be updated, replaced, or obsoleted by other documents at any time. It is inappropriate to use Internet-Drafts as reference material or to cite them other than as "work in progress."

The list of current Internet-Drafts can be accessed at <http://www.ietf.org/ietf/lid-abstracts.txt>

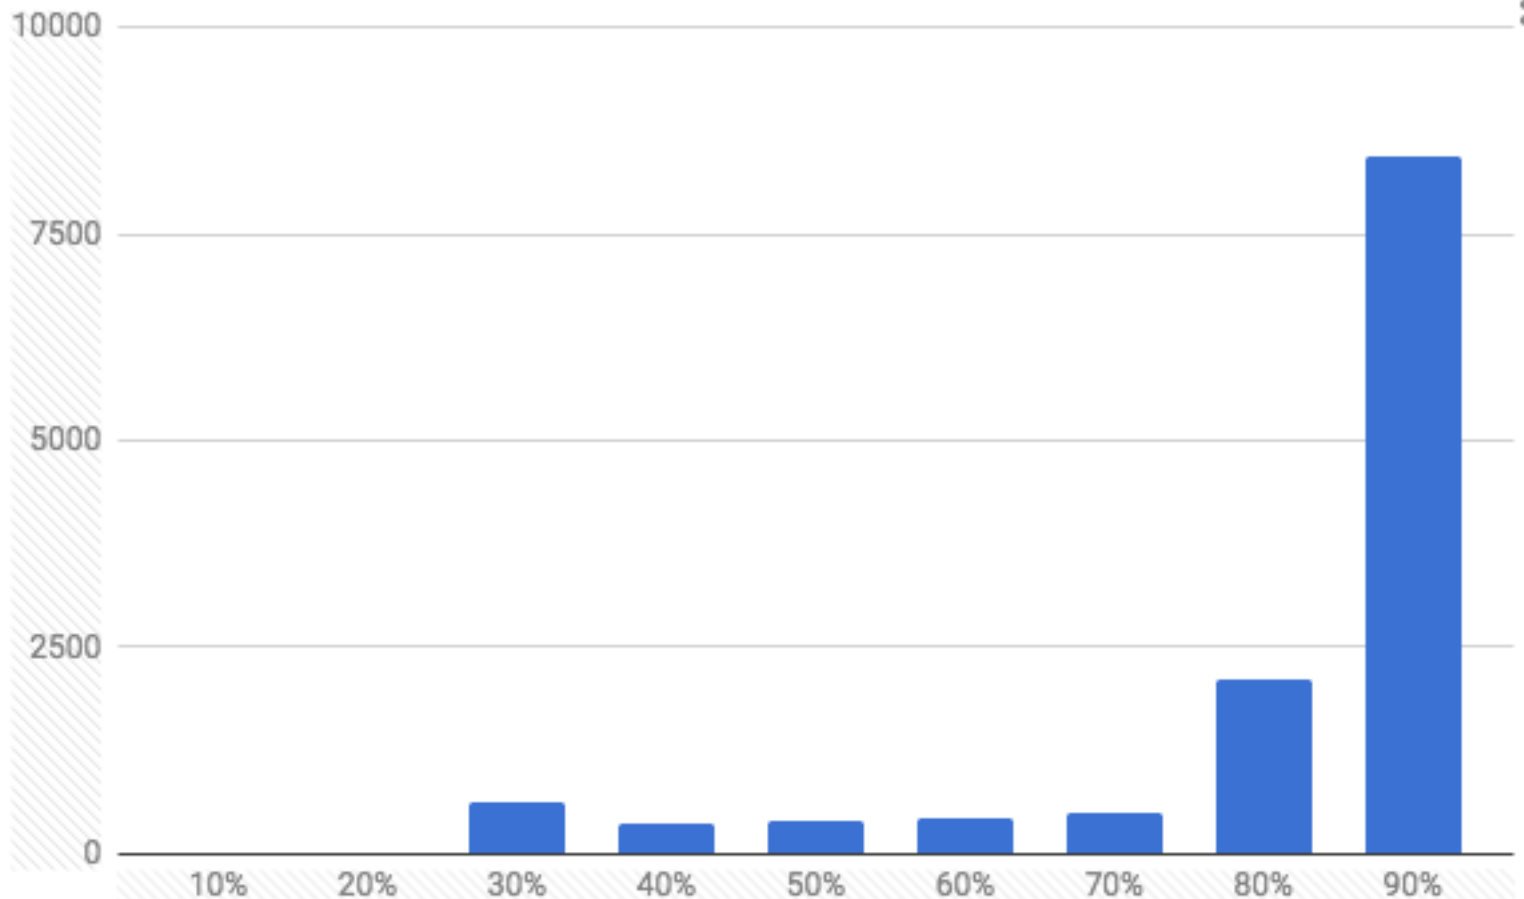
The list of Internet-Draft Shadow Directories can be accessed at <http://www.ietf.org/shadow.html>.

# Request size



source: Enabling HTTP/2 for Dropbox web services

## Compression ratio of headers <chrome://histograms>



# Clever techniques

---

- Domain sharding
- Image sprites
- JS/CSS concatenation
- Inlining





# HTTP/1.1 core problems

---

- Poor concurrency and multiplexing
- No interleaving
- Bad interaction with TCP
- Headers too large and repetitive

# HTTP/2

# RFC 7540

---

Internet Engineering Task Force (IETF)  
Request for Comments: 7540  
Category: Standards Track  
ISSN: 2070-1721

K. Belshé  
BitGo  
R. Peon  
Google, Inc  
M. Thomson, Ed.  
Mozilla  
May 2015

## Hypertext Transfer Protocol Version 2 (HTTP/2)

### Abstract

This specification describes an optimized expression of the semantics of the Hypertext Transfer Protocol (HTTP), referred to as HTTP version 2 (HTTP/2). HTTP/2 enables a more efficient use of network resources and a reduced perception of latency by introducing header field compression and allowing multiple concurrent exchanges on the same connection. It also introduces unsolicited push of representations from servers to clients.

This specification is an alternative to, but does not obsolete, the HTTP/1.1 message syntax. HTTP's existing semantics remain unchanged.

### Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 5741.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <http://www.rfc-editor.org/info/rfc7540>.

# A little history

---

- Started with Google's SPDY
- Call for proposals for HTTP/2, SPDY was used as a foundation
- Results: RFC 7540 (and RFC 7541)

# HTTP/2 design goals

---

- Solve multiplexing / interleaving
- Solve header bloat
- More efficient implementations
- Retain the semantics of HTTP/1.1: headers, caching ...



# What is HTTP/2?

## An analogy

# Sand transport

---



# Sand transport

---





# Sand transport

---



# Sand transport

---



# Sand transport



# How does it work?

# H2 components - TCP connection

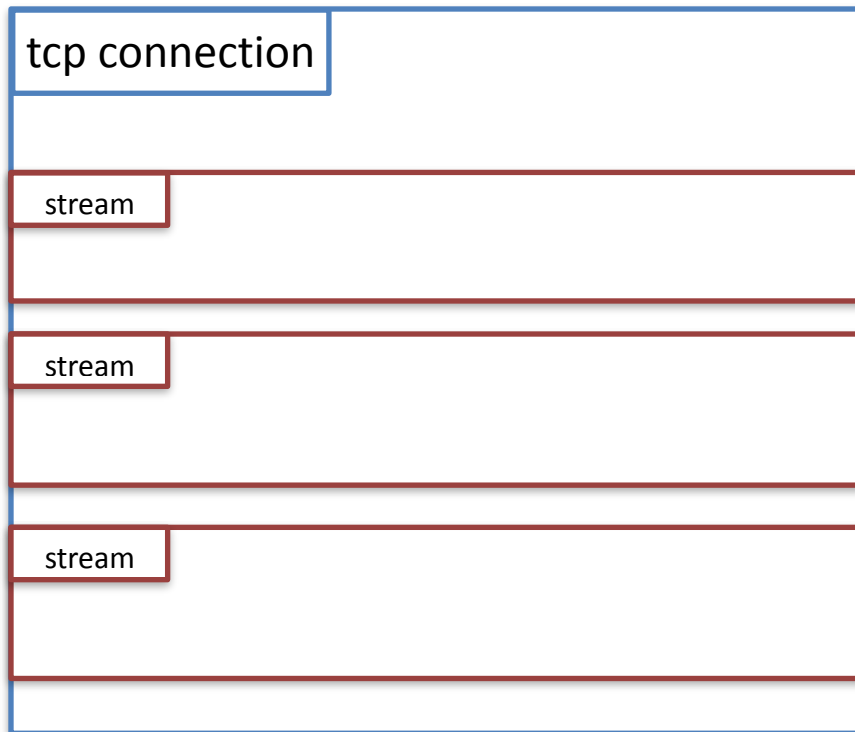
---



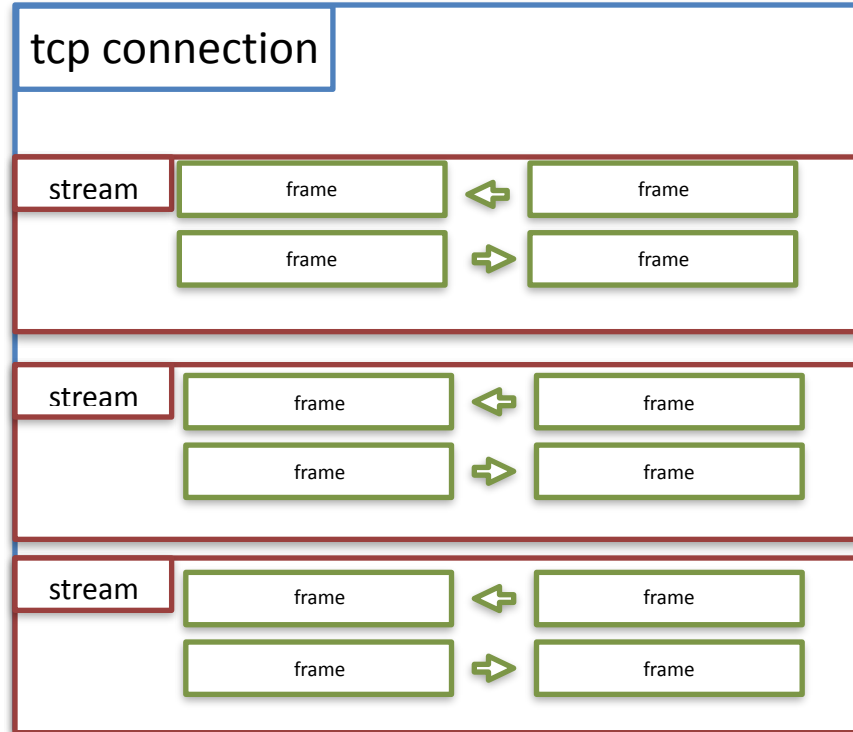
tcp connection



# H2 components - Streams



# H2 components - Frame



# Binary protocol



# HTTP/1.1 - user-agent: human

---

```
$ telnet www.fastly.com 80
Trying 151.101.25.57...
Connected to prod.www-fastly-com.map.fastlylb.net.
Escape character is '^]'.
GET / HTTP/1.1
Host: www.fastly.com
```

```
HTTP/1.1 301 Moved Permanently
Server: Varnish
Retry-After: 0
Location: https://www.fastly.com/
Content-Length: 0
Accept-Ranges: bytes
```

```
$ (printf "PRI * HTTP/2.0\r\n\r\nSM\r\n\r\n";sleep 1) | \
  openssl s_client -quiet -alpn h2 -servername www.fastly.com \
  -connect www.fastly.com:443 2> /dev/null | hexdump -C
00000000  00 00 06 04 00 00 00 00  00 00 03 00 00 00 64 00  |.....d.|
```

```
h2get> connect https://www.google.com
```

```
h2get> prefix
```

```
=====  
000: 50 52 49 20 2a 20 48 54 54 50 2f 32 2e 30 0d 0a PRI * HTTP/2.0..  
010: 0d 0a 53 4d 0d 0a 0d 0a ..SM....  
=====
```

```
h2get> settings
```

```
=====  
000: 00 00 00 04 00 00 00 00 00 .....  
=====
```

```
h2get> get
```

```
=====  
000: 00 00 23 01 05 00 00 00 01 ..#. ....  
=====
```

```
=====  
000: 82 84 87 41 8b f1 e3 c2 f3 1c f3 50 55 c8 7a 7f ...A.....PU.z.  
010: 53 03 2a 2f 2a 90 7a 8b aa 69 d2 9a c4 c0 57 08 S.*/*..z..i....W.  
020: 57 7c 1f W|.   
=====
```

```
h2get> connect https://www.google.com
```

```
h2get> prefix
```

```
=====  
000: 50 52 49 20 2a 20 48 54 54 50 2f 32 2e 30 0d 0a PRI * HTTP/2.0..  
010: 0d 0a 53 4d 0d 0a 0d 0a ..SM....  
=====
```

```
h2get> settings
```

```
=====  
000: 00 00 00 04 00 00 00 00 00 .....  
=====
```

```
h2get> get
```

```
=====  
000: 00 00 23 01 05 00 00 00 01 ..#. ....  
=====
```

```
=====  
000: 82 84 87 41 8b f1 e3 c2 f3 1c f3 50 55 c8 7a 7f ...A.....PU.z.  
010: 53 03 2a 2f 2a 90 7a 8b aa 69 d2 9a c4 c0 57 08 S.*/*.z..i....W.  
020: 57 7c 1f W|.   
=====
```

Index	Header Name	Header Value
1	:authority	
2	:method	GET
3	:method	POST
4	:path	/
5	:path	/index.html
6	:scheme	http
7	:scheme	https
8	:status	200
9	:status	204
10	:status	206
11	:status	304
12	:status	400
13	:status	404
14	:status	500
15	accept-charset	
16	accept-encoding	gzip, deflate
17	accept-language	
18	accept-ranges	
19	accept	
20	access-control-allow-origin	
21	age	
22	allow	
23	authorization	
24	cache-control	
25	content-disposition	
26	content-encoding	
27	content-language	
28	content-length	
29	content-location	
30	content-range	

# Binary protocol

---

- The good:
  - Well defined format, strict framing
  - Easier to implement
- The bad:
  - Much harder to diagnose and troubleshoot
  - Need decoding tools like Wireshark
  - Translation issues when proxying

# The connection

# A single connection

---

- Everything in HTTP/2 happens over a single, long lived connection
- Theoretically, this means better congestion management between peers
  - Standard TCP: each connection operates independently
  - Multiple connections are not competing for the same network resources



## h2 & h2c

---

- The protocol supports both secure and clear-text versions:
  - h2 (secure)
  - h2c (clear)
- Browser implementers have opted only to support the secure version

# ALPN

---

- Since we connect to port 443, we need a way to discover HTTP/2 support
- Application-Layer Protocol Negotiation — RFC 7301
- Used to be NPN, but Chrome dropped support in spring 2016

# ALPN

```
Transmission Control Protocol, Src Port: 6782 (6482), Dest Port: 443 (443), Seq: 1, Ack: 1, Len: 139
Secure Sockets Layer
  TLSv1.2 Record Layer: Handshake Protocol: Client Hello
    (packet type: handshake (22))
    Version: TLS 1.0 (0x0301)
    Length: 185
    > Handshake Protocol: Client Hello
      Handshake type: Client Hello (1)
      Length: 101
      Version: TLS 1.2 (0x0303)
      > Version
        Session ID Length: 0
        Cipher Suites Length: 34
      > Cipher Suites (12 suites)
        Compression Methods Length: 1
      > Compression Methods (1 method)
        Extensions Length: 100
      > Extension: renegotiation_info
      > Extension: server_name
      > Extension: Extended Master Secret
      > Extension: SessionTicket TLS
      > Extension: signature_algorithms
      > Extension: status_request
      > Extension: signed_certificate_timestamp
      > Extension: Application Layer Protocol Negotiation
        Type: Application Layer Protocol Negotiation (0x0030)
        Length: 14
        ALPN Extension Length: 12
        > ALPN Protocol
          alpn_string length: 2
          ALPN Next Protocol: h2
          ALPN string length: 8
          alpn_next_protocols: http/1.1
      > Extension: channel_id
      > Extension: ocsp_ext_formats
      > Extension: elliptic_curves
```

Hey!  
I support h2  
and http/1.1

# ALPN

```
Transmission Control Protocol, Src Port: 64362 (64362), Dest Port: 443 (443), Seq: 1, Ack: 1, Len: 128
Secure Sockets Layer
  TLSv1.2 Record Layer: Handshake Protocol: Client Hello
    Content Type: Handshake (22)
    Version: TLS 1.0 (0x0301)
    Length: 185
    Handshake Protocol: Client Hello
      Handshake Type: Client Hello (1)
      Length: 185
      Version: TLS 1.2 (0x0303)
      Random
      Session ID Length: 0
      Cipher Suites Length: 34
      Compression Methods Length: 1
      Extensions Length: 100
      Extension: renegotiation_info
      Extension: server_name
      Extension: Extended Master Secret
      Extension: SessionTicket TLS
      Extension: signature_algorithms
      Extension: status_request
      Extension: signed_certificate_timestamp
      Extension: Application Layer Protocol Negotiation
        Type: Application Layer Protocol Negotiation (0x0000)
        Length: 14
        ALPN Extension Length: 12
        ALPN Protocol
          alpn_string length: 2
          ALPN Next Protocol: h2
          ALPN string length: 8
          alpn_next_protocols: http/1.1
      Extension: channel_id
      Extension: ec_point_formats
      Extension: elliptic_curves
```

```
Transmission Control Protocol, Src Port: 64362 (64362), Dest Port: 64362 (64362), Seq: 1, Ack: 191, Len: 1448
Secure Sockets Layer
  TLSv1.2 Record Layer: Handshake Protocol: Server Hello
    Content Type: Handshake (22)
    Version: TLS 1.2 (0x0303)
    Length: 138
    Handshake Protocol: Server Hello
      Handshake Type: Server Hello (2)
      Length: 138
      Version: TLS 1.2 (0x0303)
      Random
      Session ID Length: 0
      Cipher Suites: TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256 (0x002F)
      Compression Methods: null (0)
      Extensions Length: 34
      Extension: server_name
      Extension: renegotiation_info
      Extension: ec_point_formats
      Extension: SessionTicket TLS
      Extension: status_request
      Extension: Application Layer Protocol Negotiation
        Type: Application Layer Protocol Negotiation (0x0000)
        Length: 8
        ALPN Extension Length: 4
        ALPN Protocol
          alpn_string length: 2
          ALPN Next Protocol: h2
```

Ok, let's do h2

# HTTP/2 and TLS

---

- MUST use SNI
- MUST disable compression
- MUST disable renegotiation
- MUST support  
TLS\_ECDHE\_RSA\_WITH\_AES\_128\_GCM\_SHA256 with the P-256 elliptic curve

# Connection re-use

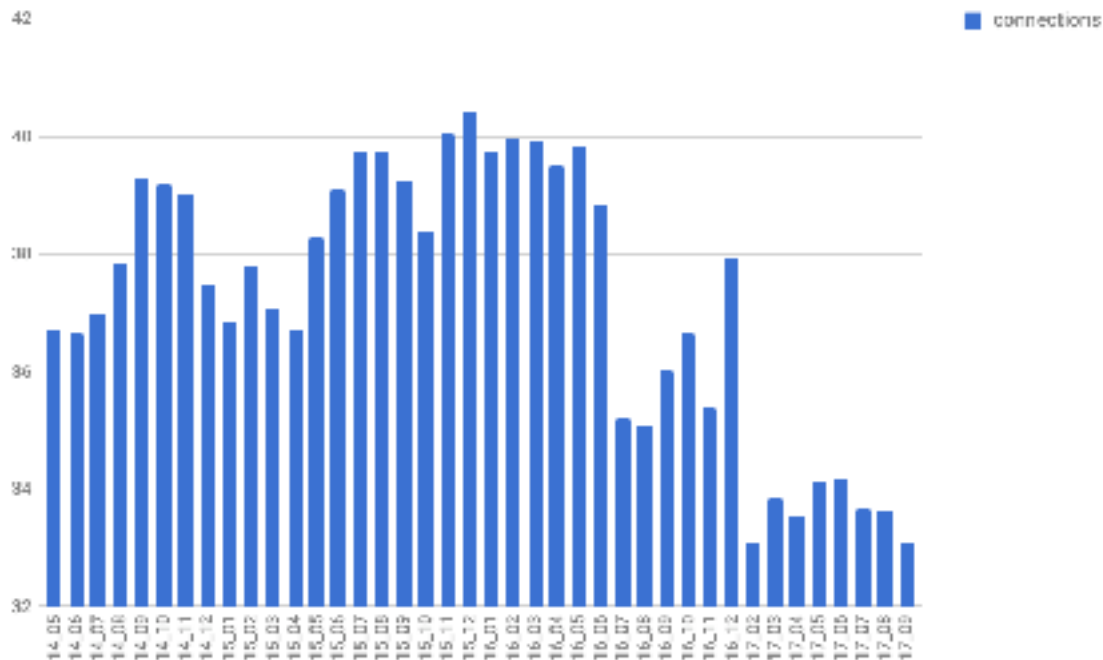
---



# Connection re-use

- Better for memory usage, server side and client side

source: [httparchive.org](http://httparchive.org)



# Connection reuse

---

- Connection can be reused between this client and another hostname if:
  - Both hostnames resolve to the same IP
  - SAN cert has both hostnames used in TLS handshake
- New status code: 421 (Misdirected request)



# Connection reuse

---

- Browsers implement this differently
  - Firefox will re-use the connection if the domains were seen using the same IP at least once
  - Chrome is more conservative and will re-use if the IPs happen to be the same



a.example.com

192.0.2.0

192.0.2.1

b.example.com

192.0.2.1

192.0.2.3



a.example.com

192.0.2.0

192.0.2.1

b.example.com

192.0.2.1

192.0.2.3



a.example.com

192.0.2.0

192.0.2.1

b.example.com

192.0.2.1

192.0.2.3

# Connection reuse

---

- Matters for testing, it's not enough to have another CNAME
- Domain sharding will behave differently

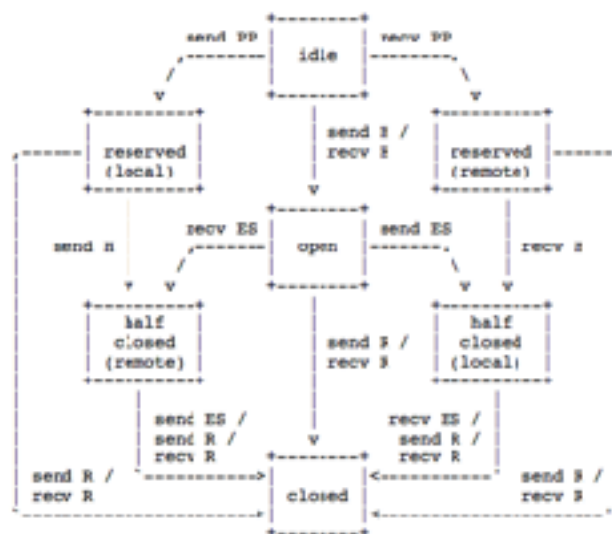
# Streams

# Streams

---

- Virtual channels of communication
  - Translate roughly to a request/response exchange
  - Client or server can initiate or terminate
- Frame order within a stream is significant
  - Frames are processed in the order they are received
- Stream IDs:
  - Client: odd, server: even, 0: connection-wide/reserved
  - Each new ID has to be larger than the previous ones
  - Stream IDs can't be reused (max is  $2^{31}$ )

# Streams



send: endpoint sends this frame  
recv: endpoint receives this frame

H: HEADERS frame (with implied CONTINUATIONS)  
PP: PUSH\_PROMISE frame (with implied CONTINUATIONS)  
ES: END\_STREAM flag  
R: RST\_STREAM frame



# Streams - <https://tools.ietf.org/html/draft-benfield-http2-debug-state-01>

---

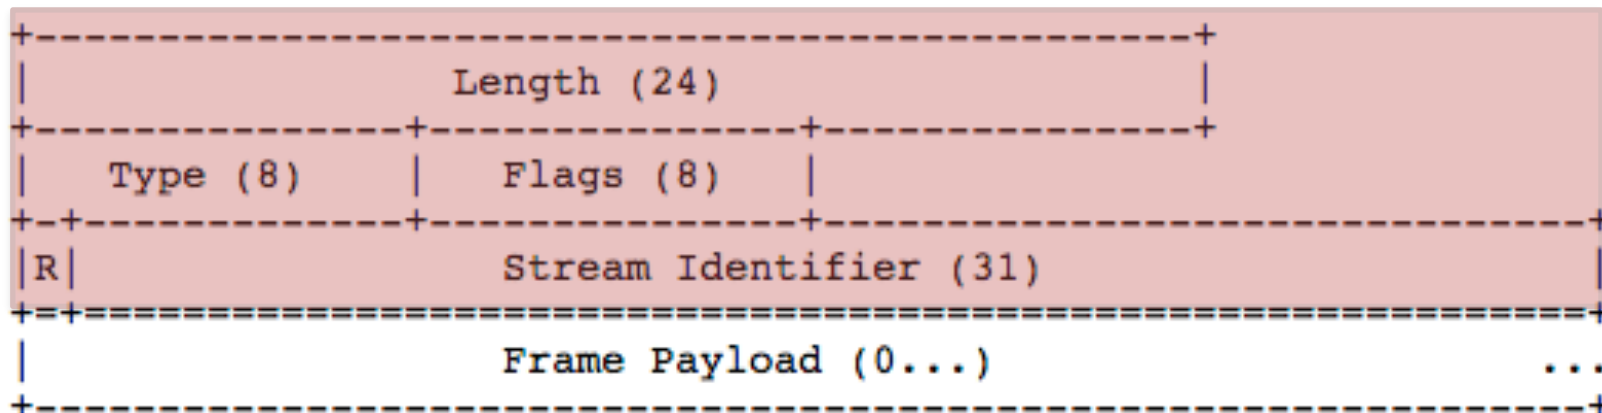
```
0 $ curl -k --http2 https://127.0.0.1:8081/.well-known/h2/state
{
  "version": "draft-01",
  "settings": {
    "SETTINGS_HEADER_TABLE_SIZE": 4096,
    "SETTINGS_ENABLE_PUSH": 0,
    "SETTINGS_MAX_CONCURRENT_STREAMS": 128,
    "SETTINGS_INITIAL_WINDOW_SIZE": 16777216,
    "SETTINGS_MAX_FRAME_SIZE": 16384
  },
  "peerSettings": {
    "SETTINGS_HEADER_TABLE_SIZE": 4096,
    "SETTINGS_ENABLE_PUSH": 1,
    "SETTINGS_MAX_CONCURRENT_STREAMS": 4294967295,
    "SETTINGS_INITIAL_WINDOW_SIZE": 65535,
    "SETTINGS_MAX_FRAME_SIZE": 16384
  },
  "connFlowIn": 65535,
  "connFlowOut": 65535,
  "streams": {
    "1": {
      "state": "HALF_CLOSED_REMOTE",
      "flowIn": 16777216,
      "flowOut": 65535,
      "dataIn": 0,
      "dataOut": 0,
      "created": 1471459379
    }
  }
}
```

# Frames

# H2 Frame

---

Fixed 9 bytes header



```
h2get> connect https://www.google.com
```

```
h2get> prefix
```

```
=====  
000: 50 52 49 20 2a 20 48 54 54 50 2f 32 2e 30 0d 0a PRI * HTTP/2.0..  
010: 0d 0a 53 4d 0d 0a 0d 0a ..SM....  
=====
```

```
h2get> settings
```

```
=====  
000: 00 00 00 04 00 00 00 00 00 .....  
=====
```

```
h2get> get
```

```
=====  
000: 00 00 23 01 05 00 00 00 01 ..#. ....  
=====
```

```
=====  
000: 82 84 87 41 8b f1 e3 c2 f3 1c f3 50 55 c8 7a 7f ...A.....PU.z.  
010: 53 03 2a 2f 2a 90 7a 8b aa 69 d2 9a c4 c0 57 08 S.*/*..z..i....W.  
020: 57 7c 1f W|.   
=====
```

# H2 Frames

---

```
GET / HTTP/1.1  
Host: www.example.com  
User-agent: Human
```

```
HTTP/1.1 200 Ok  
Server: httpserv  
Content-Type: text/html  
Content-Length: 1000
```

```
<html>  
  <body>  
    ..  
  </body>  
</html>
```

# H2 Frames

---

```
GET / HTTP/1.1
Host: www.example.com
User-agent: Human
```

Headers

```
HTTP/1.1 200 Ok
Server: httpserv
Content-Type: text/html
Content-Length: 1000
```

```
<html>
  <body>
    ..
  </body>
</html>
```

# H2 Frames

---

```
GET / HTTP/1.1  
Host: www.example.com  
User-agent: Human
```

Headers

```
HTTP/1.1 200 Ok  
Server: httpserv  
Content-Type: text/html  
Content-Length: 1000
```

Headers

```
<html>  
  <body>  
    ..  
  </body>  
</html>
```

# H2 Frames

---

```
GET / HTTP/1.1
Host: www.example.com
User-agent: Human
```

Headers

```
HTTP/1.1 200 Ok
Server: httpserv
Content-Type: text/html
Content-Length: 1000
```

Headers

```
<html>
  <body>
    ..
  </body>
</html>
```

Data

Data

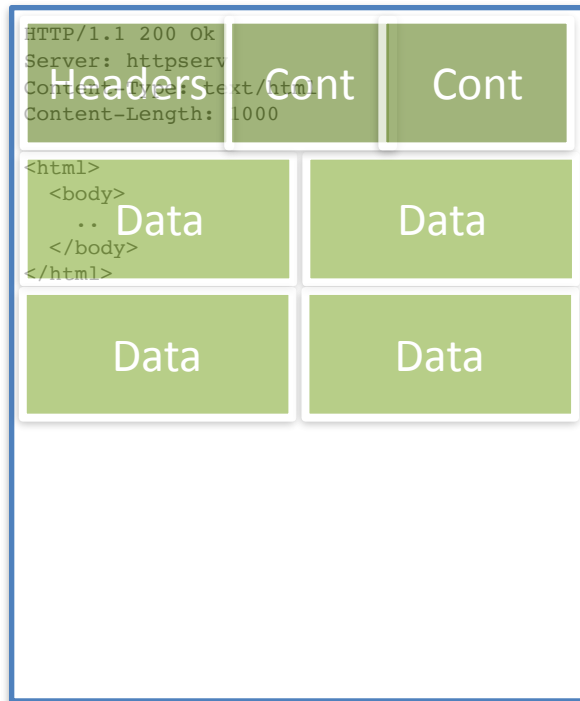
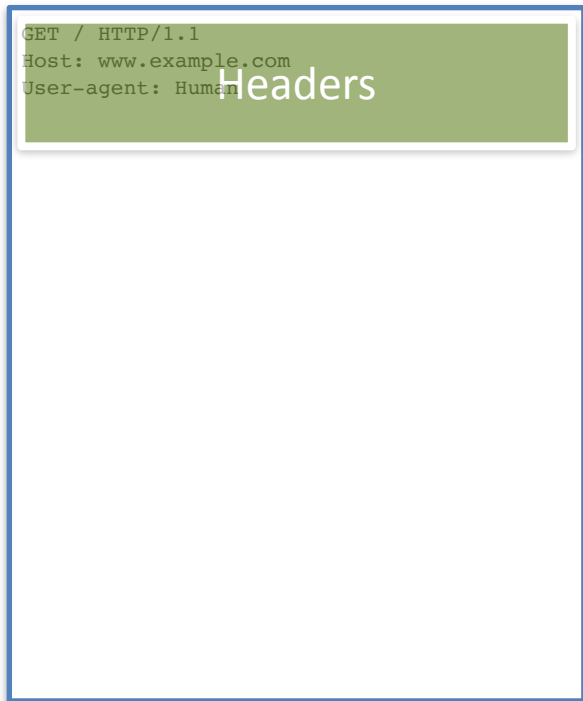
Data

Data



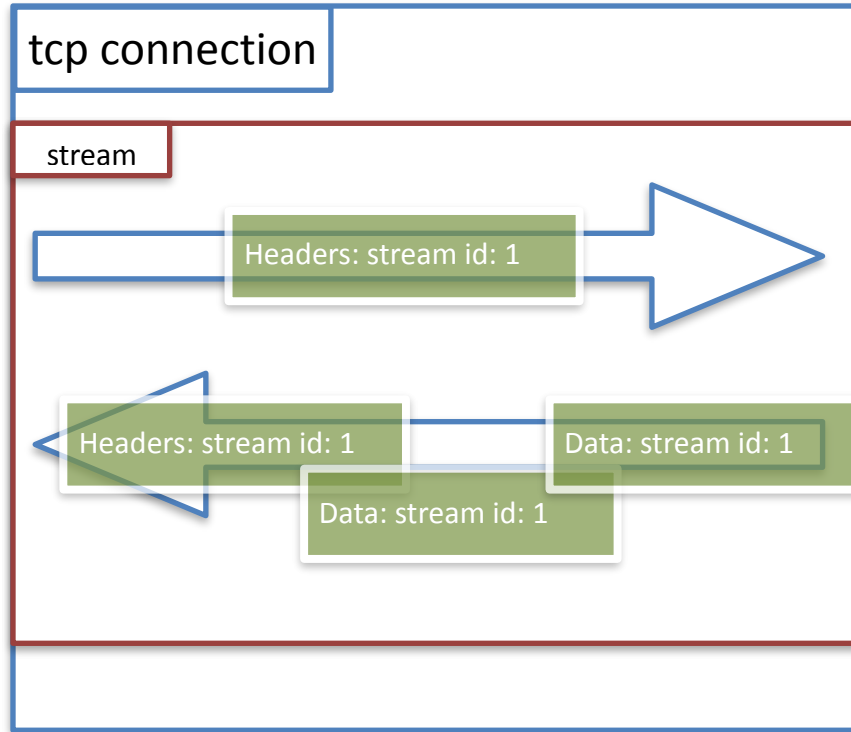
# H2 Frames

---

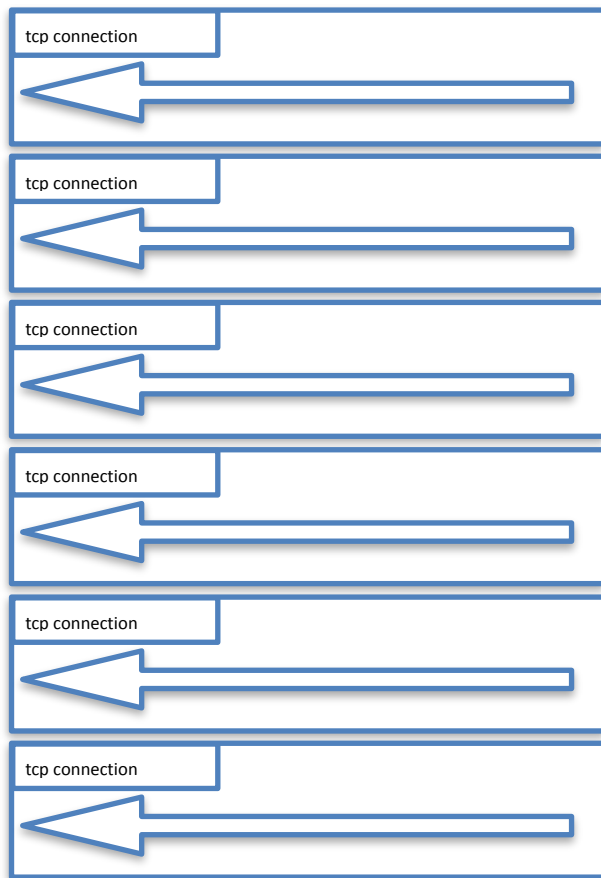


# Protocol flow

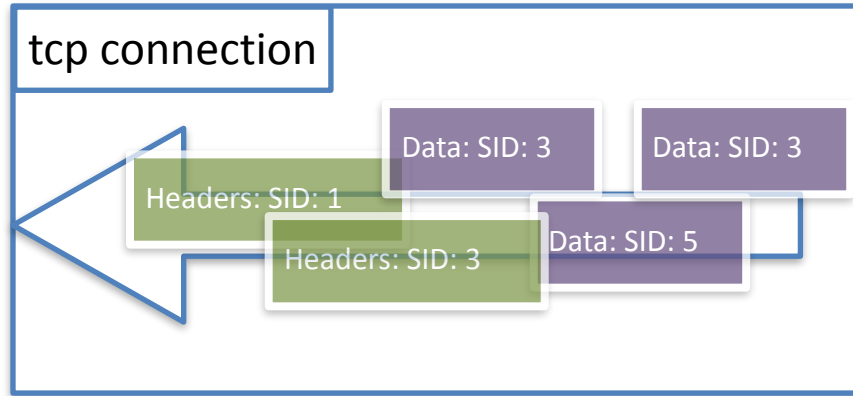
# H2 Frames



# HTTP/1.1 used to be

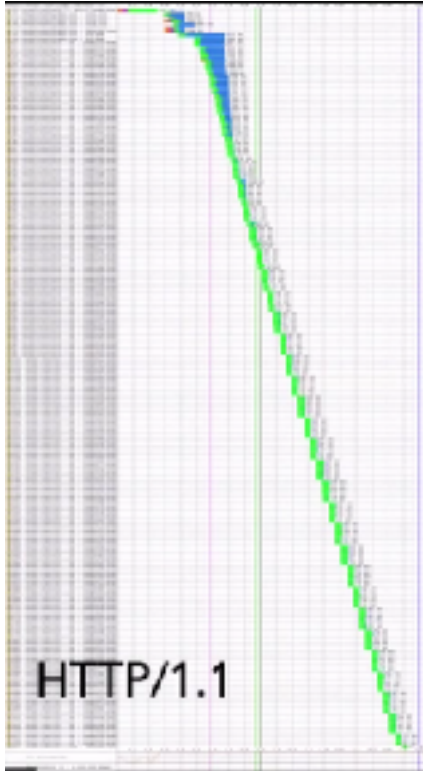


# H2 Frames



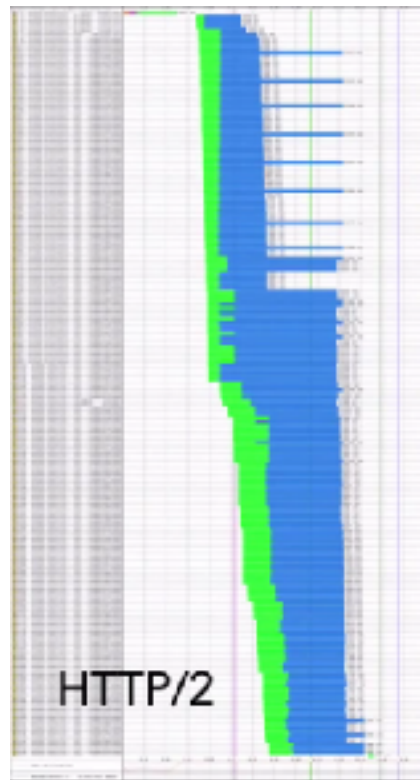
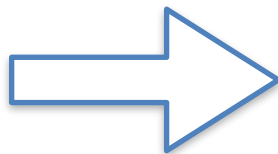
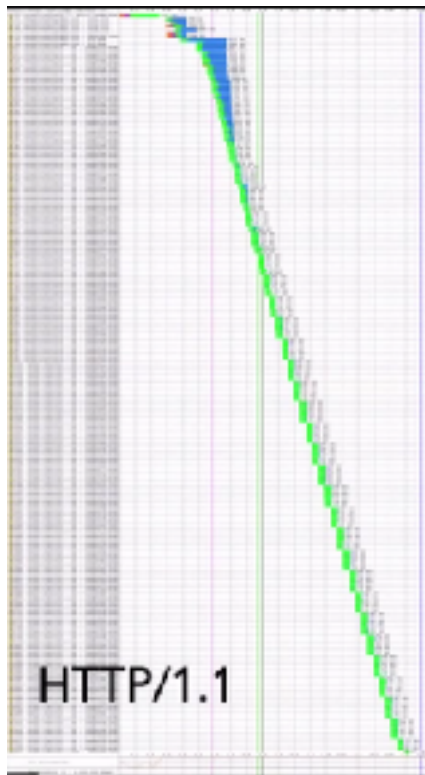
# H2 Frames

---



# H2 Frames

---



# H2 Frames - Wireshark

57	0.000000	10.100.8.96	151.101.53.57	TCP	78	56221→443 [SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=1024 TSval=1378205703 TSecr=0 SACK_PERM=1
58	0.021017	151.101.53.57	10.100.8.96	TCP	74	443→56221 [SYN, ACK] Seq=0 Ack=1 Win=28960 Len=0 MSS=1460 SACK_PERM=1 TSval=1691902767 TSecr=1378205703 WS=512
59	0.000072	10.100.8.96	151.101.53.57	TCP	66	56221→443 [ACK] Seq=1 Ack=1 Win=4194304 Len=0 TSval=1378205724 TSecr=1691902767
60	0.000335	10.100.8.96	151.101.53.57	TLSv1..	269	Client Hello
61	0.022025	151.101.53.57	10.100.8.96	TLSv1..	1514	Server Hello
62	0.000002	151.101.53.57	10.100.8.96	TCP	1514	[TCP segment of a reassembled PDU]
63	0.000001	151.101.53.57	10.100.8.96	ILP	1514	[ILP segment of a reassembled PDU]
64	0.000002	151.101.53.57	10.100.8.96	TLSv1..	912	CertificateStatus, Server Key Exchange, Server Hello Done
65	0.000063	10.100.8.96	151.101.53.57	TCP	66	56221→443 [ACK] Seq=204 Ack=2397 Win=4191232 Len=0 TSval=1378205746 TSecr=1691902773
66	0.000003	10.100.8.96	151.101.53.57	TCP	66	56221→443 [ACK] Seq=204 Ack=3191 Win=4189184 Len=0 TSval=1378205746 TSecr=1691902773
67	0.000123	10.100.8.96	151.101.53.57	TCP	66	[TCP Window Update] 56221→443 [ACK] Seq=204 Ack=3191 Win=4193200 Len=0 TSval=1378205740 TSecr=1691902773
68	0.002540	10.100.8.96	151.101.53.57	TLSv1..	192	Client Key Exchange, Change Cipher Spec, Finished
69	0.000244	10.100.8.96	151.101.53.57	HTTP2	229	Magic, SETTINGS, WINDOW_UPDATE, PRIORITY, PRIORITY, PRIORITY, PRIORITY, PRIORITY
70	0.000003	10.100.8.96	151.101.53.57	HTTP2	331	HEADERS, WINDOW_UPDATE
71	0.025011	151.101.53.57	10.100.8.96	TCP	66	443→56221 [ACK] Seq=3191 Ack=708 Win=31200 Len=0 TSval=1691902779 TSecr=1378205740
72	0.000002	151.101.53.57	10.100.8.96	TLSv1..	308	New Session Ticket, Change Cipher Spec, Finished
73	0.000079	10.100.8.96	151.101.53.57	TCP	66	56221→443 [ACK] Seq=758 Ack=3433 Win=4193200 Len=0 TSval=1378205759 TSecr=1691902779
74	0.000191	151.101.53.57	10.100.8.96	HTTP2	125	SETTINGS, SETTINGS
75	0.000030	10.100.8.96	151.101.53.57	TCP	66	56221→443 [ACK] Seq=758 Ack=3432 Win=4193200 Len=0 TSval=1378205759 TSecr=1691902779
76	0.000065	10.100.8.96	151.101.53.57	HTTP2	104	SETTINGS
77	0.000256	151.101.53.57	10.100.8.96	TCP	1514	[TCP segment of a reassembled PDU]
78	0.000002	151.101.53.57	10.100.8.96	HTTP2	1210	HEADERS, DATA
79	0.000024	10.100.8.96	151.101.53.57	TCP	66	56221→443 [ACK] Seq=796 Ack=8004 Win=4191232 Len=0 TSval=1378205759 TSecr=1691902779
80	0.000005	151.101.53.57	10.100.8.96	TCP	1514	[TCP segment of a reassembled PDU]
81	0.000002	151.101.53.57	10.100.8.96	TLSv1..	1406	[SSL segment of a reassembled PDU]
82	0.000002	151.101.53.57	10.100.8.96	TCP	1514	[TCP segment of a reassembled PDU]
83	0.000001	151.101.53.57	10.100.8.96	TLSv1..	1514	[SSL segment of a reassembled PDU]
84	0.000001	151.101.53.57	10.100.8.96	TCP	1514	[TCP segment of a reassembled PDU]
85	0.000002	151.101.53.57	10.100.8.96	TLSv1..	1298	[SSL segment of a reassembled PDU]



# H2 Frames - Wireshark

74	0.000191	151.101.53.57	10.100.8.96	HTTP2	125	SETTINGS, SETTINGS
75	0.000230	10.100.8.96	151.101.53.57	TCP	66	56221→443 [ACK] Seq=758 Ack=5492 Win=4193280 Len=0 TSval=1378285769 TSecr=1691982779
76	0.000265	10.100.8.96	151.101.53.57	HTTP2	104	SETTINGS
77	0.000256	151.101.53.57	10.100.8.96	TCP	1514	[TCP segment of a reassembled PDU]
78	0.000202	151.101.53.57	10.100.8.96	HTTP2	1210	HEADERS, DATA
79	0.000224	10.100.8.96	151.101.53.57	TCP	66	56221→443 [ACK] Seq=796 Ack=8084 Win=4193280 Len=0 TSval=1378285769 TSecr=1691982779
80	0.000285	151.101.53.57	10.100.8.96	TCP	1514	[TCP segment of a reassembled PDU]
81	0.000282	151.101.53.57	10.100.8.96	TLSv1	1406	[SSL segment of a reassembled PDU]
82	0.000282	151.101.53.57	10.100.8.96	TCP	1514	[TCP segment of a reassembled PDU]

▶ Frame 78: 1210 bytes on wire (9680 bits), 1210 bytes captured (9680 bits) on interface 0

▶ Ethernet II, Src: Juniper1\_95:c9:f8 (44:f4:77:95:c9:f8), Dst: Apple\_d3:60:55 (0c:4d:e8:d3:60:55)

▶ Internet Protocol Version 4, Src: 151.101.53.57, Dst: 10.100.8.96

▶ Transmission Control Protocol, Src Port: 443, Dst Port: 56221, Seq: 6948, Ack: 758, Len: 1144

▶ [2 Reassembled TCP Segments (2592 bytes): #77(1448), #78(1144)]

▼ Secure Sockets Layer

▼ TLSv1.2 Record Layer: Application Data Protocol: http2

Content Type: Application Data (23)

Version: TLS 1.2 (0x0303)

Length: 2587

Encrypted Application Data: 0b95ae7ddcfe44d9571681ba225a329233705a57c896419e...

▼ HyperText Transfer Protocol 2

▶ Stream: HEADERS, Stream ID: 13, Length 428

▶ Stream: DATA, Stream ID: 13, Length 2117

# H2 Frames - Wireshark

75	0.000000	192.168.0.96	255.255.255.255	TCP	8856221443 [ACK] Seq=758 Ack=5492 Win=4191232 Len=0 TSval=1378285769 TSecr=3891822719
76	0.000000	192.168.0.96	255.255.255.255	HTTP2	104 SETTINGS
77	0.000000	155.181.53.57	30.160.3.96	TCP	1534 [TCP segment of a reassembled PDU]
78	0.000000	155.181.53.57	30.160.3.96	HTTP2	1230 HEADERS, DATA
79	0.000000	192.168.0.96	255.255.255.255	TCP	8856221443 [ACK] Seq=758 Ack=8854 Win=4191232 Len=0 TSval=1378285769 TSecr=3891822719
80	0.000000	155.181.53.57	30.160.3.96	TCP	1534 [TCP segment of a reassembled PDU]
81	0.000000	155.181.53.57	30.160.3.96	HTTP2	1406 SSL segment of a reassembled PDU
82	0.000000	155.181.53.57	30.160.3.96	TCP	1534 [TCP segment of a reassembled PDU]
83	0.000000	155.181.53.57	30.160.3.96	HTTP2	1534 [SSL segment of a reassembled PDU] [TCP segment of a reassembled PDU]
84	0.000000	155.181.53.57	30.160.3.96	TCP	1534 [TCP segment of a reassembled PDU]
85	0.000000	155.181.53.57	30.160.3.96	HTTP2	1708 [SSL segment of a reassembled PDU]
86	0.000000	192.168.0.96	255.255.255.255	TCP	8856221443 [ACK] Seq=758 Ack=12872 Win=4191232 Len=0 TSval=1378285769 TSecr=3891822719

Length: 1587  
Encrypted Application Data: 8295ae705c7e44c05f7883a225b329213705a57c358418c...

HyperText Transfer Protocol 2

Stream: HEADERS, Stream ID: 13, Length: 428  
Length: 428  
Type: HEADERS (1)  
Flags: 0x00  
E... .. = Reserved: 0x00  
... .. = Stream Identifier: 13  
[Pad Length: 0]  
Header Block Fragments: 885144d724b75a049ac0ff104906a099a8d4a03e208b...  
[Header Length: 341]  
[Header Count: 21]  
Header: :status: 200  
Header: date: Thu, 10 Nov 2016 23:42:40 GMT  
Header: content-encoding: gzip  
Header: content-language: en  
Header: content-type: text/html; charset=utf-8  
Header: link: <https://www.fastly.com/> ;rel="canonical", <https://www.fastly.com/> ;rel="shortlink"  
Header: via: 1.1 varnish  
Header: x-content-type-options: nosniff  
Header: last-modified: Thu, 10 Nov 2016 13:24:13 GMT  
Header: via: 1.1 varnish  
Header: fastly-debug-digest: 2ba610b18f437ab61a96211ed1a3a82df3a76dfe9097c154a75b1d27217  
Header: access-control-allow-headers: DRYDEN  
Header: via: 1.1 varnish  
Header: x-served-by: cache-sjc3132-514, cache-sjc3132-514  
Header: x-cache: HIT, HIT  
Header: x-cache-hit: 2, 2007  
Header: x-timer: S1475821806.129250.750.000  
Header: vary: Cookie, Fastly-reqid, accept-encoding  
Header: cache-control: max-age=3600, private, must-revalidate

:status: 200

# Frame types

# Frame types

---

Frame Type	Code	Section
DATA	0x0	Section 6.1
HEADERS	0x1	Section 6.2
PRIORITY	0x2	Section 6.3
RST_STREAM	0x3	Section 6.4
SETTINGS	0x4	Section 6.5
PUSH_PROMISE	0x5	Section 6.6
PING	0x6	Section 6.7
GOAWAY	0x7	Section 6.8
WINDOW_UPDATE	0x8	Section 6.9
CONTINUATION	0x9	Section 6.10

# Frame types

---

Frame Type	Code	Section
DATA	0x0	Section 6.1
HEADERS	0x1	Section 6.2
PRIORITY	0x2	Section 6.3
RST_STREAM	0x3	Section 6.4
SETTINGS	0x4	Section 6.5
PUSH_PROMISE	0x5	Section 6.6
PING	0x6	Section 6.7
GOAWAY	0x7	Section 6.8
WINDOW_UPDATE	0x8	Section 6.9
CONTINUATION	0x9	Section 6.10

# Frame types

---

Frame Type	Code	Section
DATA	0x0	Section 6.1
HEADERS	0x1	Section 6.2
PRIORITY	0x2	Section 6.3
RST_STREAM	0x3	Section 6.4
<b>SETTINGS</b>	<b>0x4</b>	<b>Section 6.5</b>
PUSH_PROMISE	0x5	Section 6.6
PING	0x6	Section 6.7
GOAWAY	0x7	Section 6.8
WINDOW_UPDATE	0x8	Section 6.9
CONTINUATION	0x9	Section 6.10

# Frame types

---

Frame Type	Code	Section
DATA	0x0	Section 6.1
HEADERS	0x1	Section 6.2
PRIORITY	0x2	Section 6.3
<b>RST_STREAM</b>	<b>0x3</b>	<b>Section 6.4</b>
SETTINGS	0x4	Section 6.5
PUSH_PROMISE	0x5	Section 6.6
PING	0x6	Section 6.7
<b>GOAWAY</b>	<b>0x7</b>	<b>Section 6.8</b>
WINDOW_UPDATE	0x8	Section 6.9
CONTINUATION	0x9	Section 6.10

# Frame types

---

Frame Type	Code	Section
DATA	0x0	Section 6.1
HEADERS	0x1	Section 6.2
PRIORITY	0x2	Section 6.3
RST_STREAM	0x3	Section 6.4
SETTINGS	0x4	Section 6.5
PUSH_PROMISE	0x5	Section 6.6
<b>PING</b>	<b>0x6</b>	<b>Section 6.7</b>
GOAWAY	0x7	Section 6.8
WINDOW_UPDATE	0x8	Section 6.9
CONTINUATION	0x9	Section 6.10



# Frame types

---

Frame Type	Code	Section
DATA	0x0	Section 6.1
HEADERS	0x1	Section 6.2
PRIORITY	0x2	Section 6.3
RST_STREAM	0x3	Section 6.4
SETTINGS	0x4	Section 6.5
PUSH_PROMISE	0x5	Section 6.6
PING	0x6	Section 6.7
GOAWAY	0x7	Section 6.8
WINDOW_UPDATE	0x8	Section 6.9
CONTINUATION	0x9	Section 6.10

# Frame types

---

- Extensible
  - Currently discussed new frames:
    - ORIGIN HTTP/2 Frame
    - CACHE\_DIGEST Frame
    - H2 shared compression dictionaries
- Implementations should ignore unknown frames

# SETTINGS Frame

---

- Defines parameters for the connection
  - Max concurrent streams
  - Max window size
  - ...
- Are sent and acknowledged by both sides
- Can be sent at any time

# SETTINGS frame

No.	Time	Source	Destination	Protocol	Length	Info
6036	26.871477	192.168.1.161	151.101.41.63	HTTP2	223	Magic, SETTINGS, WINDOW_UPDATE, PRIORITY, PRIORITY, PRIORITY, PRIORITY, PRIORITY
6037	26.871507	192.168.1.161	151.101.41.63	HTTP2	227	HEADERS, WINDOW_UPDATE
6038	26.884571	151.101.41.63	192.168.1.161	HTTP2	125	SETTINGS, SETTINGS
6040	26.884906	192.168.1.161	151.101.41.63	HTTP2	104	SETTINGS
6056	27.176189	151.101.41.63	192.168.1.161	HTTP2	1015	HEADERS, DATA
6066	27.465413	192.168.1.161	94.46.159.28	HTTP2	211	Magic, SETTINGS, WINDOW_UPDATE, PRIORITY, PRIORITY, PRIORITY, PRIORITY, PRIORITY
6067	27.465500	192.168.1.161	94.46.159.28	HTTP2	254	HEADERS, WINDOW_UPDATE
6088	27.470278	192.168.1.161	94.46.159.28	HTTP2	92	SETTINGS
6089	27.479843	94.46.159.28	192.168.1.161	HTTP2	92	SETTINGS
6119	27.507205	94.46.159.28	192.168.1.161	HTTP2	1514	HEADERS, DATA
6136	27.531281	94.46.159.28	192.168.1.161	HTTP2	1514	DATA
6139	27.532658	94.46.159.28	192.168.1.161	HTTP2	496	DATA
6148	27.714003	192.168.1.161	151.101.41.63	HTTP2	145	HEADERS, WINDOW_UPDATE
6154	28.022941	151.101.41.63	192.168.1.161	HTTP2	941	HEADERS, DATA
6156	28.023766	192.168.1.161	151.101.41.63	HTTP2	140	HEADERS, WINDOW_UPDATE
6164	28.350468	151.101.41.63	192.168.1.161	HTTP2	917	HEADERS, DATA

# SETTINGS frame

---

- ▶ Internet Protocol Version 4, Src: 192.168.1.161, Dst: 151.101.41.63
- ▶ Transmission Control Protocol, Src Port: 58977 (58977), Dst Port: 443 (443), Seq: 318, Ack: 6566, Len: 157
- ▶ Secure Sockets Layer
- ▼ HyperText Transfer Protocol 2
  - ▶ Stream: Magic
  - ▶ Stream: SETTINGS, Stream ID: 0, Length 12
  - ▶ Stream: WINDOW\_UPDATE, Stream ID: 0, Length 4
  - ▶ Stream: PRIORITY, Stream ID: 3, Length 5
  - ▶ Stream: PRIORITY, Stream ID: 5, Length 5
  - ▶ Stream: PRIORITY, Stream ID: 7, Length 5
  - ▶ Stream: PRIORITY, Stream ID: 9, Length 5
  - ▶ Stream: PRIORITY, Stream ID: 11, Length 5

# SETTINGS Frame

```
▼ HyperText Transfer Protocol 2
  ▼ Stream: Magic
    Magic: PRI * HTTP/2.0\r\n\r\nSM\r\n\r\n
  ▼ Stream: SETTINGS, Stream ID: 0, Length 12
    Length: 12
    Type: SETTINGS (4)
    ▶ Flags: 0x00
      0... .... = Reserved: 0x00000000
      .000 0000 0000 0000 0000 0000 0000 0000 - Stream Identifier: 0
    ▼ Settings - Initial Windows size : 131072
      Settings Identifier: Initial Windows size (4)
      Initial Windows Size: 131072
    ▼ Settings - Max frame size : 16384
      Settings Identifier: Max frame size (5)
      Max frame size: 16384
    ▶ Stream: WINDOW UPDATE, Stream ID: 0, Length 4
    ▶ Stream: PRIORITY, Stream ID: 3, Length 5
    ▶ Stream: PRIORITY, Stream ID: 5, Length 5
    ▶ Stream: PRIORITY, Stream ID: 7, Length 5
    ▶ Stream: PRIORITY, Stream ID: 9, Length 5
    ▶ Stream: PRIORITY, Stream ID: 11, Length 5
```

# SETTINGS Frame

---

- SETTINGS\_HEADER\_TABLE\_SIZE 4K
- SETTINGS\_ENABLE\_PUSH true
- SETTINGS\_MAX\_CONCURRENT\_STREAMS unlimited
- SETTINGS\_INITIAL\_WINDOW\_SIZE 64K
- SETTINGS\_MAX\_FRAME\_SIZE 16K
- SETTINGS\_MAX\_HEADER\_LIST\_SIZE unlimited

# HPACK (RFC 7541)



# HPACK

55 pages



- Addresses the issue of header bloat

```
$ cat req_bytes1 | awk '{print $4}' | grep -oE "[0-9]+"  
| awk '{ print int(($0-1)/1000)*1000 + 1000; }' | head  
-n 1000 | sort -n | uniq -c | sort -n  
      6 5000  
     18 4000  
     89 3000  
    192 2000  
    695 1000
```

# HPACK

---

- More requests will fit in the initial window
- Smaller headers can help save a roundtrip

# HPACK

---

- Addresses the issue of header bloat
- Two ways:
  - Huffman encoding
  - Indexed tables

# HPACK - Huffman encoding

---

'u' ( 85)	1110000	70 [ 7]
'v' ( 86)	1110001	71 [ 7]
'w' ( 87)	1110010	72 [ 7]
'x' ( 88)	1111100	fc [ 8]
'y' ( 89)	1110011	73 [ 7]
'z' ( 90)	1111101	fd [ 8]
'[' ( 91)	1111111 11011	1ffb [13]
'\' ( 92)	1111111 1111110 000	7fff0 [19]
' ' ( 93)	1111111 11100	1ffn [13]
'^' ( 94)	1111111 111100	3ffn [14]
'_' ( 95)	100010	22 [ 6]
'`' ( 96)	1111111 1111101	7ffd [15]
'a' ( 97)	00011	3 [ 5]
'b' ( 98)	100011	23 [ 6]
'c' ( 99)	00100	4 [ 5]
'd' (100)	100100	24 [ 6]
'e' (101)	00101	5 [ 5]
'f' (102)	100101	25 [ 6]
'g' (103)	100110	26 [ 6]
'h' (104)	100111	27 [ 6]
'i' (105)	00110	6 [ 5]
'j' (106)	1110100	74 [ 7]
'k' (107)	1110101	75 [ 7]
'l' (108)	101000	28 [ 6]
'm' (109)	101001	29 [ 6]

# HPACK - static table

---

Index	Header Name	Header Value
1	:authority	
2	:method	GET
3	:method	POST
4	:path	/
5	:path	/index.html
6	:scheme	http
7	:scheme	https
8	:status	200
9	:status	204
10	:status	206
11	:status	304
12	:status	400
13	:status	404
14	:status	500
15	accept-charset	
16	accept-encoding	gzip, deflate
17	accept-language	
18	accept-ranges	
19	accept	
20	access-control-allow-origin	
21	age	
22	allow	
23	authorization	
24	cache-control	
25	content-disposition	
26	content-encoding	
27	content-language	
28	content-length	
29	content-location	
30	content-range	

# HPACK - dynamic tables

---

- Dynamic tables size is dictated via the SETTINGS frame
- The compression context is the connection, longer lived connections are favoured

# HPACK - dynamic table

```
:method: GET
:scheme: http
:path: /
:authority: www.example.com

Hex dump of encoded data:

8286 8441 8cf1 e3c2 e5f2 3a8b a0ab 90f4 | ...A.....:k....
ff                                         | .

Decoding process:

82                                         | == Indexed - Add ==
                                         |   idx = 2
                                         | -> :method: GET
86                                         | == Indexed - Add ==
                                         |   idx = 6
                                         | -> :scheme: http
84                                         | == Indexed - Add ==
                                         |   idx = 4
                                         | -> :path: /
41                                         | == Literal indexed ==
                                         |   Indexed name (idx = 1)
                                         |     :authority
8c                                         |   Literal value (len = 12)
                                         |     Huffman encoded:
f1e3 c2e5 f23a 63a0 ab90 f4ff           | .....:k.....
                                         |   Decoded:
                                         |   www.example.com
                                         | -> :authority:
                                         |   www.example.com

Dynamic Table (after decoding):

[ 1] {s = 57} :authority: www.example.com
    Table size: 57
```



# HPACK - dynamic table

```
:method: GET
:scheme: http
:path: /
:authority: www.example.com

Hex dump of encoded data:
8286 8441 8c91 e3c2 e5f2 3a8b a0ab 90f4 | ...A.....:k....
ff                                     | .

Decoding process:
82 == Indexed - Add ==
   idx = 2
86 -> :method: GET
   == Indexed - Add ==
   idx = 6
84 -> :scheme: http
   == Indexed - Add ==
   idx = 4
41 -> :path: /
   == Literal indexed ==
   Indexed name (idx = 1)
   :authority
   Literal value (len = 12)
   Huffman encoded:
8c .....:k.....
   Decoded:
   www.example.com
   -> :authority:
   www.example.com

Dynamic Table (after decoding):
[ 1] {s = 57} :authority: www.example.com
    Table size: 57
```

# HPACK - dynamic table

```
:method: GET
:scheme: http
:path: /
:authority: www.example.com
:cache-control: no-cache

Hex dump of encoded data:
8?86 848e f886 a8eb 1064 8cbf | ....X....d..

Decoding process:
8? | == Indexed - Add ==
86 |   idx = 2
84 | -> :method: GET
84 | == Indexed - Add ==
84 |   idx = 6
84 | -> :scheme: http
84 | == Indexed - Add ==
84 |   idx = 4
84 | -> :path: /
84 | == Indexed - Add ==
84 |   idx = 82
84 | -> :authority:
84 |   www.example.com
58 | == Literal indexed ==
86 |   Indexed name (idx = 24)
86 |     :cache-control
86 |     Literal value (len = 8)
86 |     Huffman encoded:
a8eb 1064 8cbf | ...d..
               | Decoded:
               | no-cache
               | -> :cache-control: no-cache

Dynamic Table (after decoding):
[ 1] {s = 53} :cache-control: no-cache
[ 2] {s = 57} :authority: www.example.com
Table size: 110
```

# HPACK - security aspects

---

- Because compression can reveal information (<https://en.wikipedia.org/wiki/CRIME>), you can opt out compression.
- It's possible to never compress information that's sent, but you can't control what the peer will do

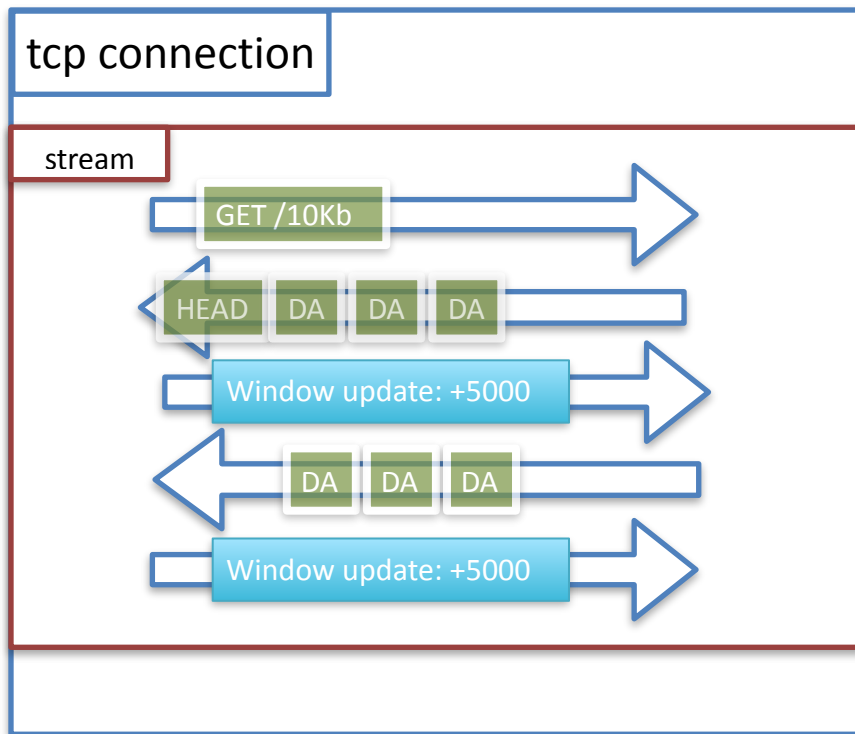
# Flow control

# Flow control basics

---

- Goal: provide a mechanism for each side to control the flow of data - on top of TCP
- Hop by hop mechanism
  - client -> proxy -> server: 2 flow controls
- Credit system controlled by the receiver

# Flow control



# Flow control

---

- Per connection and per stream
- Connection always defaults to 64k
- Each stream defaults to settings
  - but that can be overridden by SETTINGS, INITIAL\_WINDOW\_SIZE
- WINDOW\_UPDATE is used to make it bigger
- Max is  $2^{31}$  bytes

# Flow control

---

- Cannot be disabled
- Only applies to DATA frames
- Negative windows are legal, they just mean the peer must not send any more DATA frames



# Priorities

# Prioritization basics

---

- Tools help manage HTTP/2's concurrency
- RFC doesn't specify the implementation details
- Can be set using HEADERS or PRIORITY frames
- PRIORITY can be sent for streams in any state

# Prioritization basics

---

- Weight
  - 0-255 prioritizes siblings
- Parent stream id
  - Allows to build a dependency tree
- Exclusive flag

# Flow control - Example 1

---

- A gets 3/4 of the resources
- B gets 1/4 of the resources

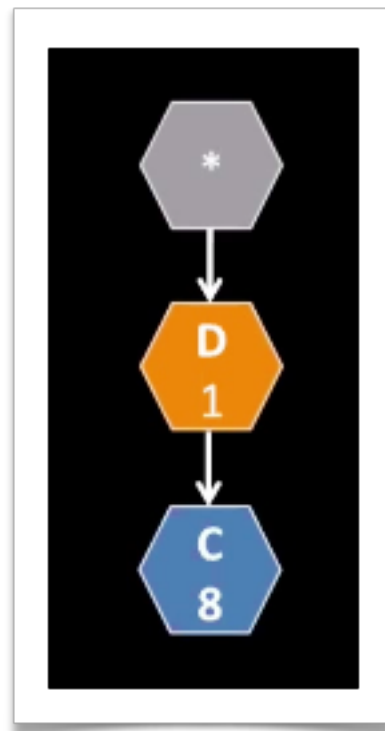


Example by Ilya Grigorik

# Flow control - Example 2

---

- D gets all resources
- After D is done, C gets all resources
- Weights have no meanings since there are no siblings



Example by Ilya Grigorik

# Flow control - Example 3

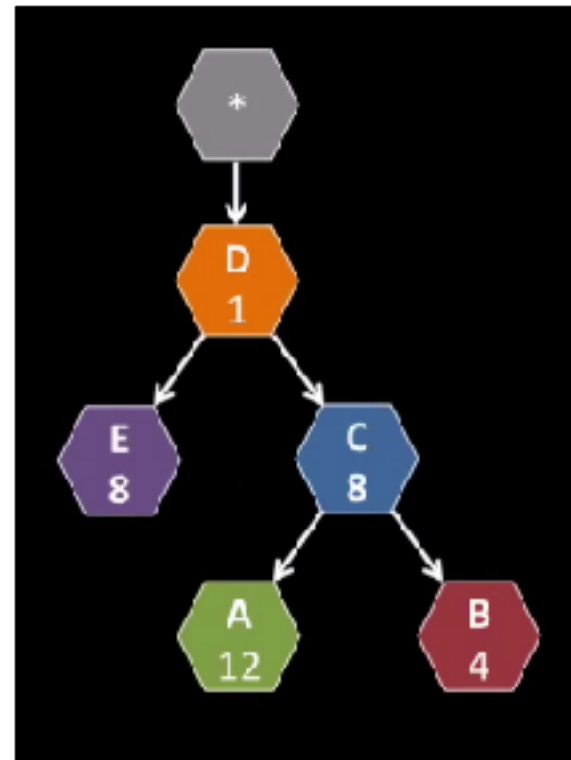
- D gets all resources
- After D is done, C gets all resources
- After C is done A gets 3/4 and B gets 1/4



Example by Ilya Grigorik

# Flow control - Example 4

- D gets all resources
- After D is done, E and C get half the resources
- After C is done A gets 3/4 and B gets 1/4



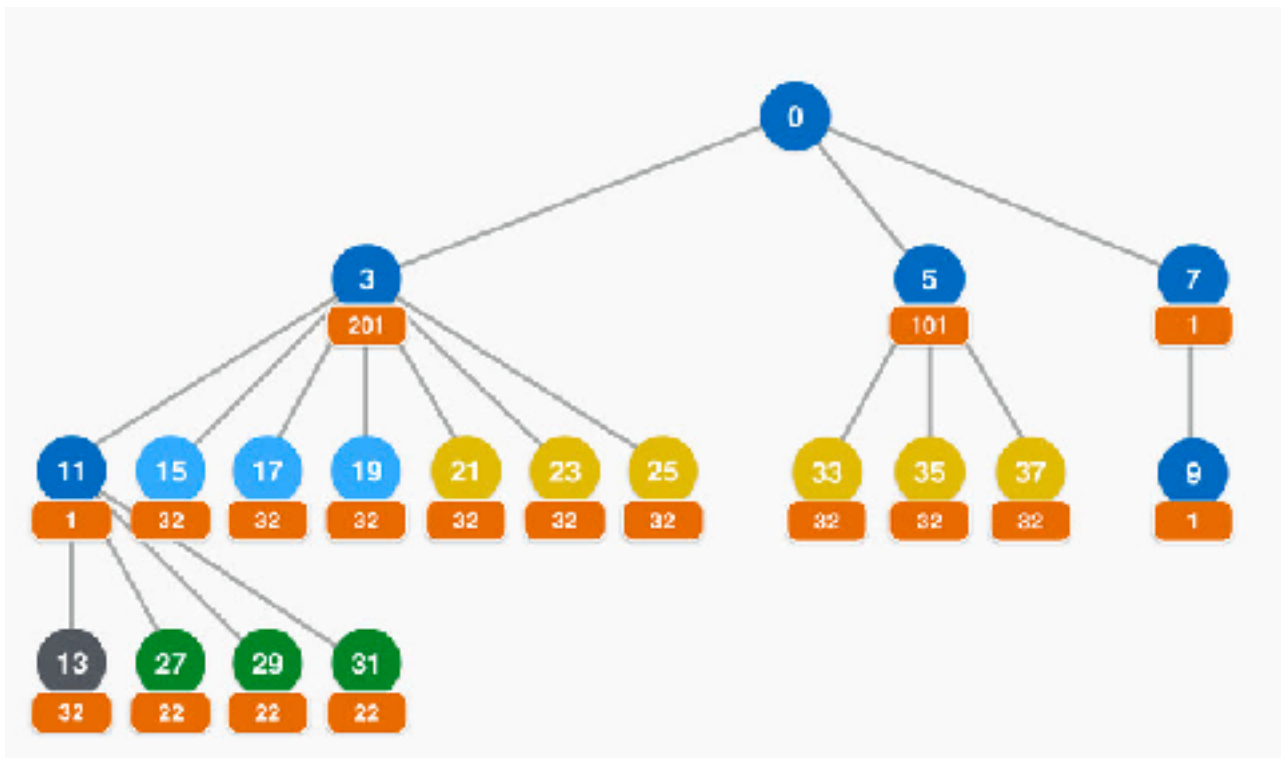
Example by Ilya Grigorik

# Browser trees

Moto Ishizawa from Yahoo! Japan



# Firefox priority tree



# SETTINGS frame

No.	Time	Source	Destination	Protocol	Length	Info
6036	26.871477	192.168.1.161	151.101.41.63	HTTP2	223	Magic, SETTINGS, WINDOW_UPDATE, PRIORITY, PRIORITY, PRIORITY, PRIORITY, PRIORITY
6037	26.871507	192.168.1.161	151.101.41.63	HTTP2	227	HEADERS, WINDOW_UPDATE
6038	26.884571	151.101.41.63	192.168.1.161	HTTP2	125	SETTINGS, SETTINGS
6040	26.884906	192.168.1.161	151.101.41.63	HTTP2	104	SETTINGS
6056	27.176189	151.101.41.63	192.168.1.161	HTTP2	1015	HEADERS, DATA
6066	27.465413	192.168.1.161	94.46.159.28	HTTP2	211	Magic, SETTINGS, WINDOW_UPDATE, PRIORITY, PRIORITY, PRIORITY, PRIORITY, PRIORITY
6087	27.465500	192.168.1.161	94.46.159.28	HTTP2	254	HEADERS, WINDOW_UPDATE
6088	27.470278	192.168.1.161	94.46.159.28	HTTP2	92	SETTINGS
6089	27.479843	94.46.159.28	192.168.1.161	HTTP2	92	SETTINGS
6119	27.507205	94.46.159.28	192.168.1.161	HTTP2	1514	HEADERS, DATA
6136	27.531281	94.46.159.28	192.168.1.161	HTTP2	1514	DATA
6139	27.532658	94.46.159.28	192.168.1.161	HTTP2	496	DATA
6148	27.714003	192.168.1.161	151.101.41.63	HTTP2	145	HEADERS, WINDOW_UPDATE
6154	28.022941	151.101.41.63	192.168.1.161	HTTP2	941	HEADERS, DATA
6156	28.023766	192.168.1.161	151.101.41.63	HTTP2	140	HEADERS, WINDOW_UPDATE
6164	28.350468	151.101.41.63	192.168.1.161	HTTP2	917	HEADERS, DATA

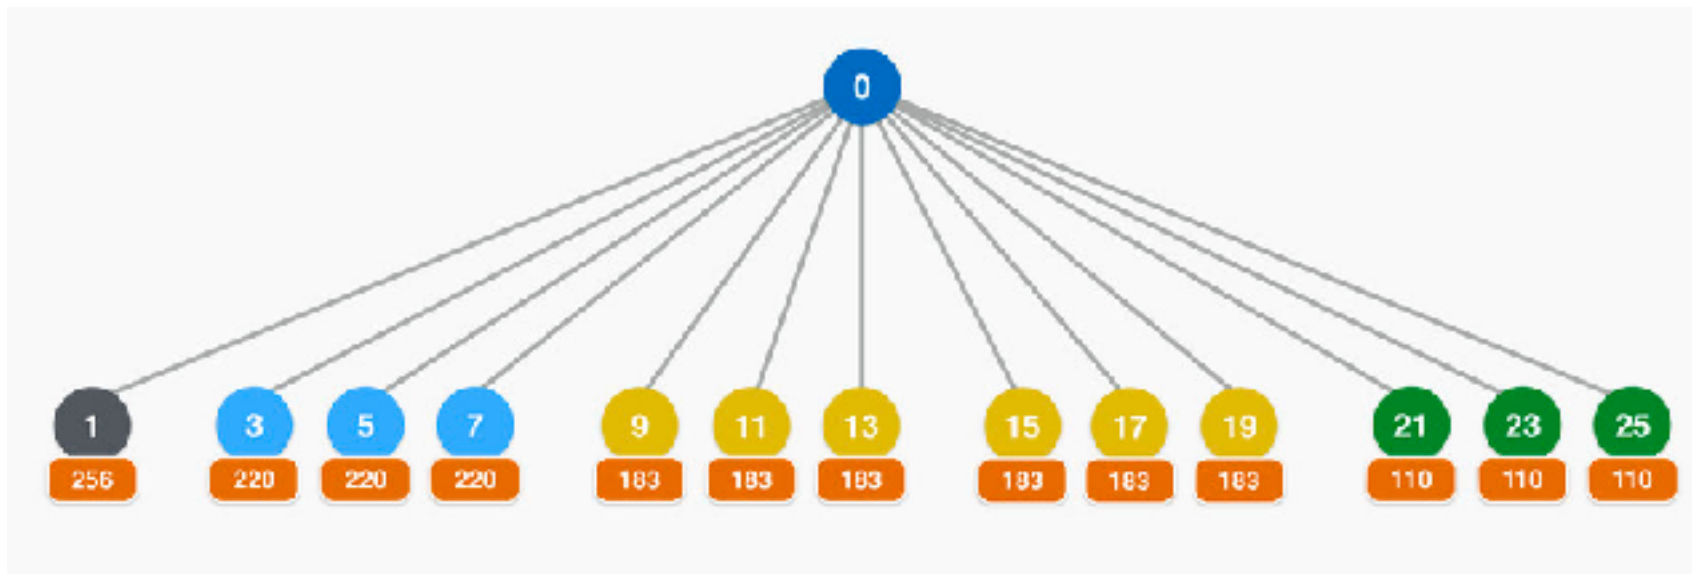
# SETTINGS frame

---

- ▶ Internet Protocol Version 4, Src: 192.168.1.161, Dst: 151.101.41.63
- ▶ Transmission Control Protocol, Src Port: 58977 (58977), Dst Port: 443 (443), Seq: 318, Ack: 6566, Len: 157
- ▶ Secure Sockets Layer
- ▼ HyperText Transfer Protocol 2
  - ▶ Stream: Magic
  - ▶ Stream: SETTINGS, Stream ID: 0, Length 12
  - ▶ Stream: WINDOW\_UPDATE, Stream ID: 0, Length 4
  - ▶ Stream: PRIORITY, Stream ID: 3, Length 5
  - ▶ Stream: PRIORITY, Stream ID: 5, Length 5
  - ▶ Stream: PRIORITY, Stream ID: 7, Length 5
  - ▶ Stream: PRIORITY, Stream ID: 9, Length 5
  - ▶ Stream: PRIORITY, Stream ID: 11, Length 5

# Chrome priority tree

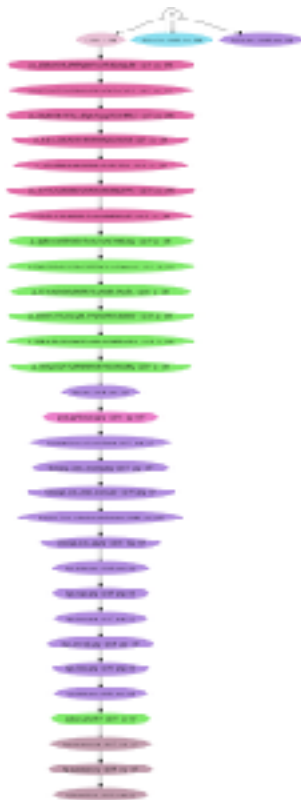
---



Understanding HTTP/2 prioritization by Moto Ishizawa

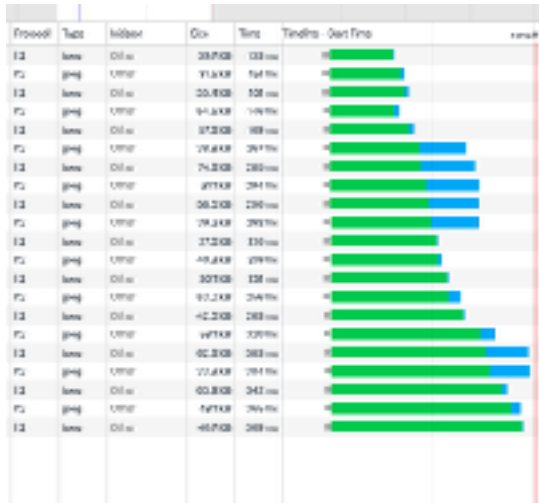
# Chrome priority tree - new

---



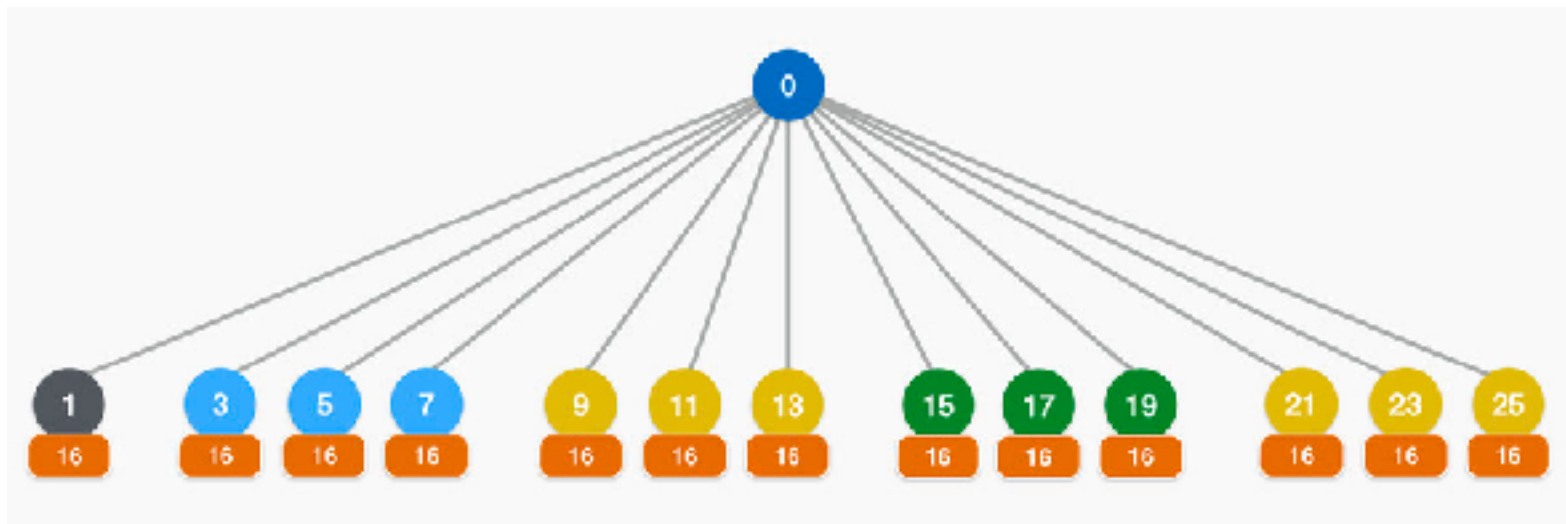
# Chrome priority tree - new

- All requests use exclusive dependencies
- Waterfalls might appear similar to http/1
- Will have HOL blocking for big objects



# MS Edge priority tree

---



Understanding HTTP/2 prioritization by Moto Ishizawa

# Server push



# Server push basics

---

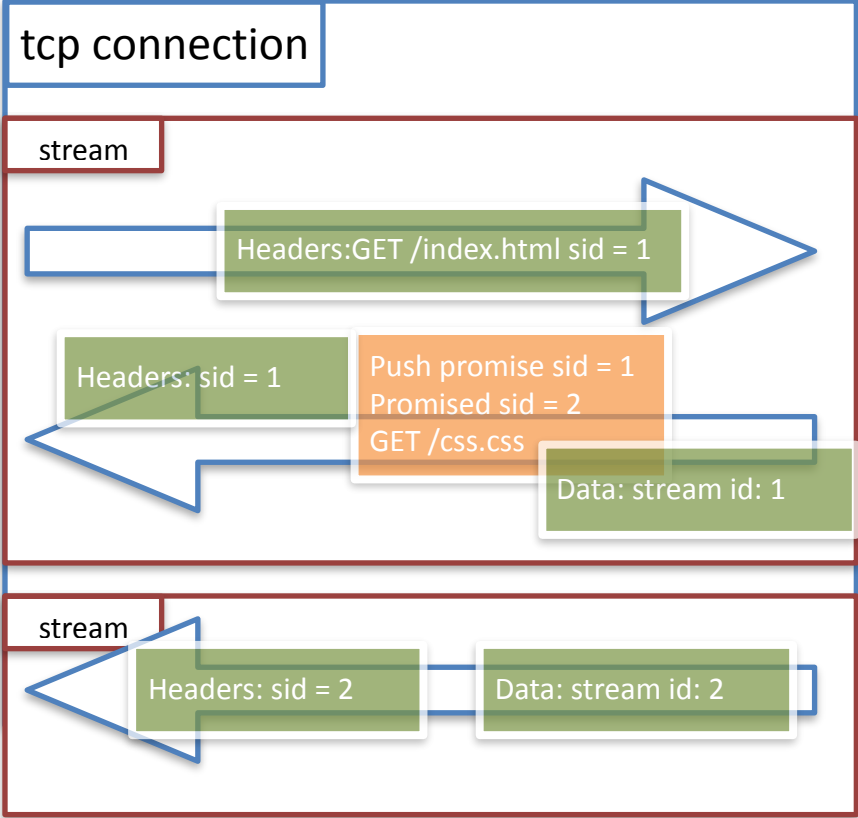
- Ability to push a resource before the client requests it
  - Possibly before the client even knows it needs it
- Operates hop by hop

# Server push basics

---

- Only a server can push
- The capability to accept pushes is advertised through the SETTINGS frame

# Push promise



# PUSH\_PROMISE frame

---

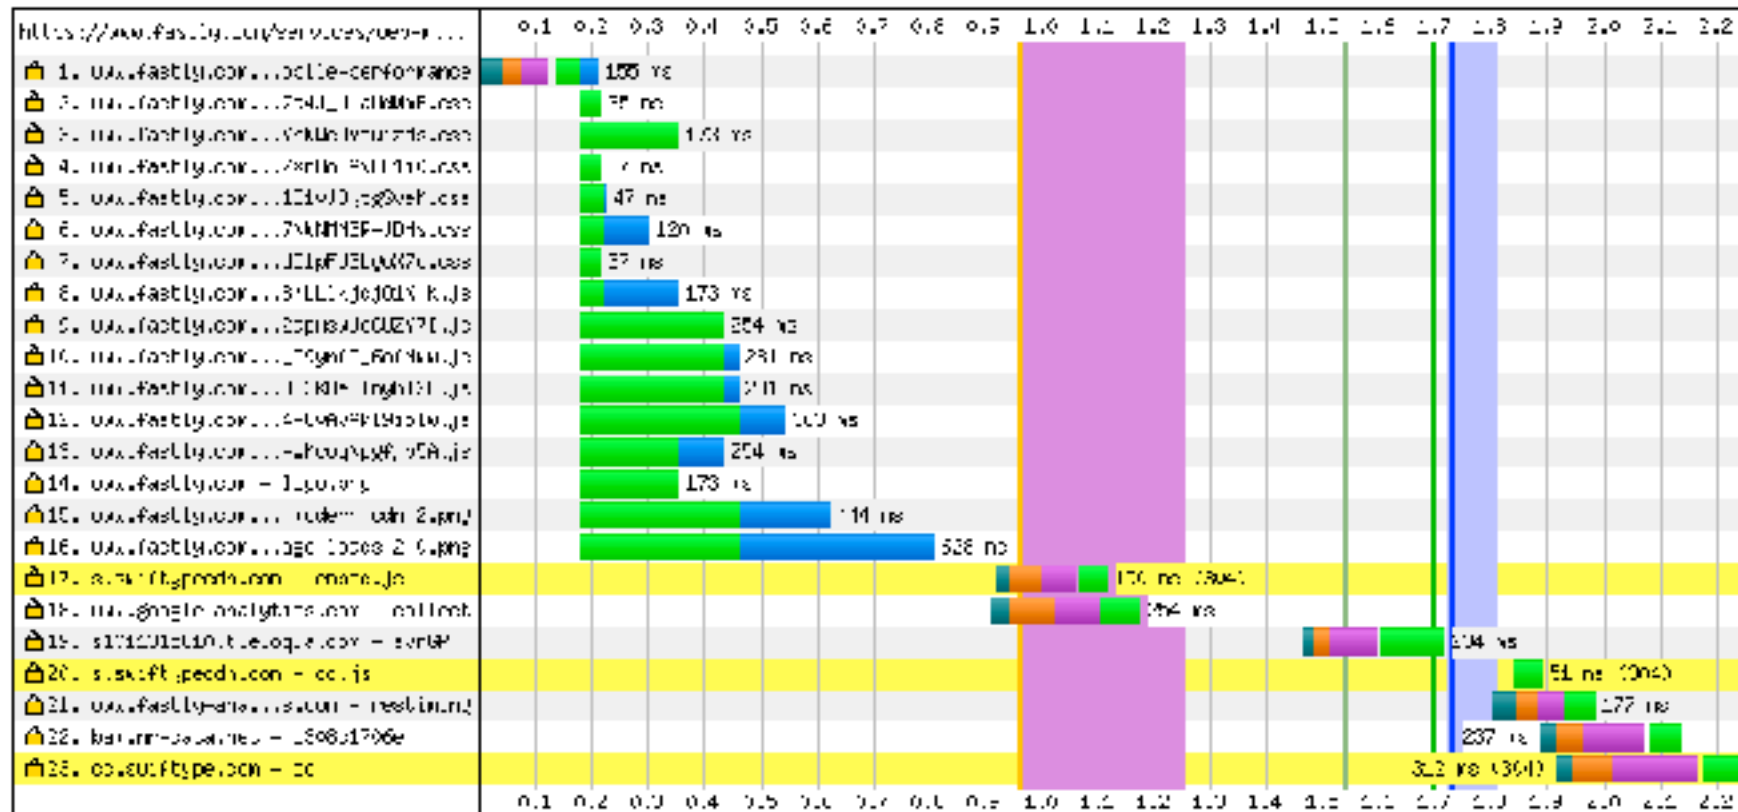
- Key elements:
  - Would-be request headers for the resource (can be followed by CONTINUATION frames)
  - Promised stream id (even)
- Sent over a client initiated stream
- Client can reject with RST\_STREAM

# Server push

---

- What can be pushed:
  - Safe method (GET/HEAD/TRACE/OPTIONS)
  - No request body
- Client can control how many pushed streams it gets through the `MAX_CONCURRENT_STREAMS` entry in `SETTINGS`

# Server push in action



# Getting to know the client's cache

---

- HTTP/2 extension:
  - <http://httpwg.org/http-extensions/cache-digest.txt>
- Cookie based: H2O's CASPer (cache-aware server push)
  - [https://h2o.example.net/configure/http2\\_directives.html#http2-casper](https://h2o.example.net/configure/http2_directives.html#http2-casper)

# Google's Rules of Thumb for HTTP/2 Push

---

- Push just enough resources to fill idle network time, and no more
- Push resources in evaluation-dependence order
- Consider using special strategies to track the client-side cache - see previous slide
- Use the right cookies when pushing resources that vary by cookie
- Use server push to fill the initial cwnd and consider using preload links to reveal the remaining critical or hidden resources
- <https://docs.google.com/document/d/1K0NykTXBbbbTlv60t5MyJvXjqKGsCVNYHyLEXIYmV0/edit>



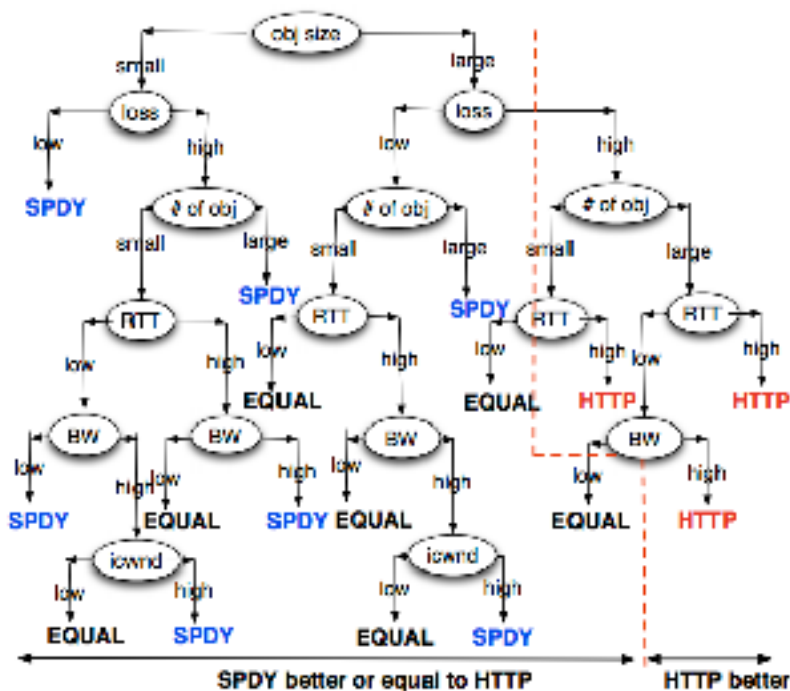
# Two clear uses cases

---

- Use server thinking time
  - 1xx early hints
  - 307 trick with H2O
- Push from the edge

# Performance

(it's complicated)



How Speedy is SPDY? - Xiao Sophia Wang, Aruna Balasubramanian, Arvind Krishnamurthy, and David Wetherall, University of Washington

# [webpagetest.org](https://webpagetest.org)

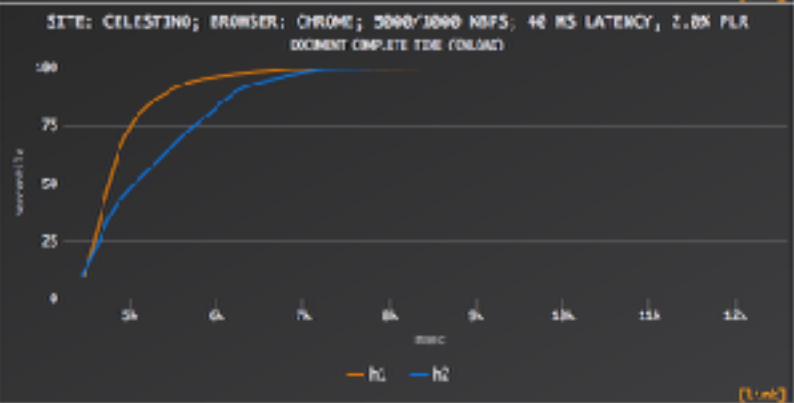
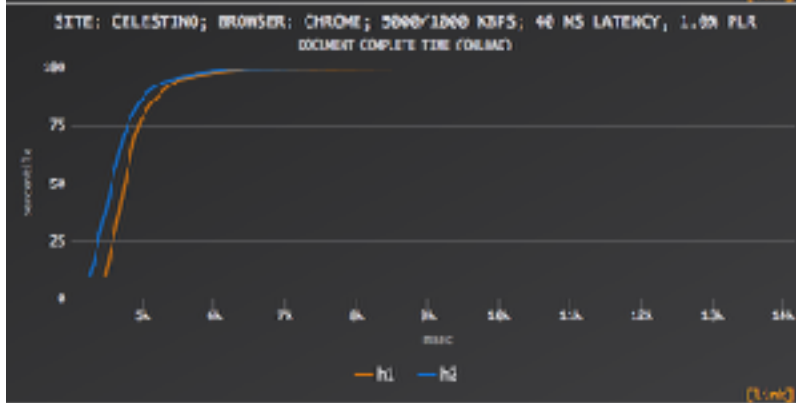
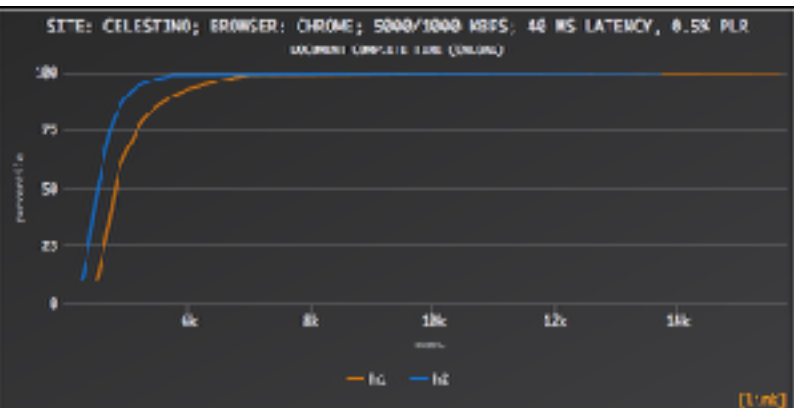
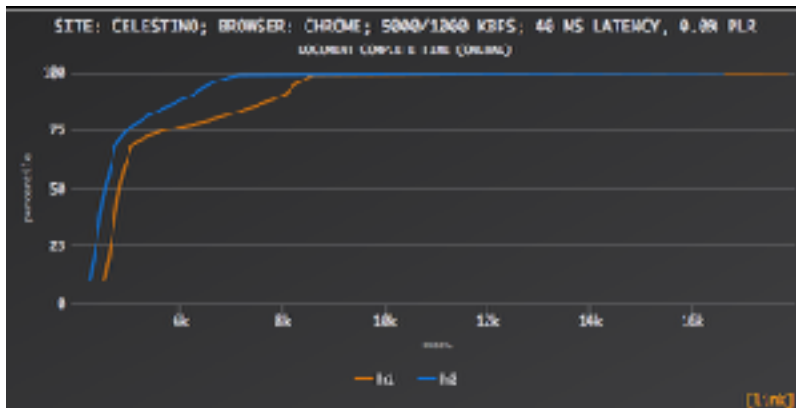
---

- Scriptable
- Performance tests for varying latencies, bandwidths, browsers
- Records videos, load times for all requests, tcpdump captures
- Can be self hosted: <https://github.com/WPO-Foundation/webpagetest>

# “Classic” web page

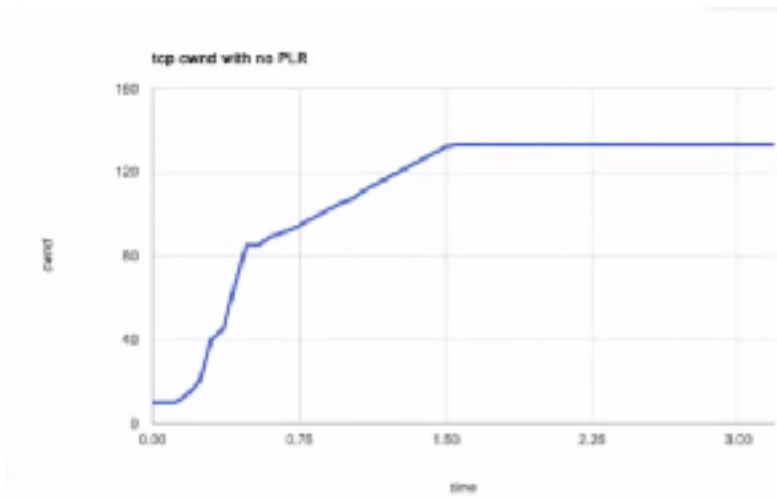


# “Classic” web page



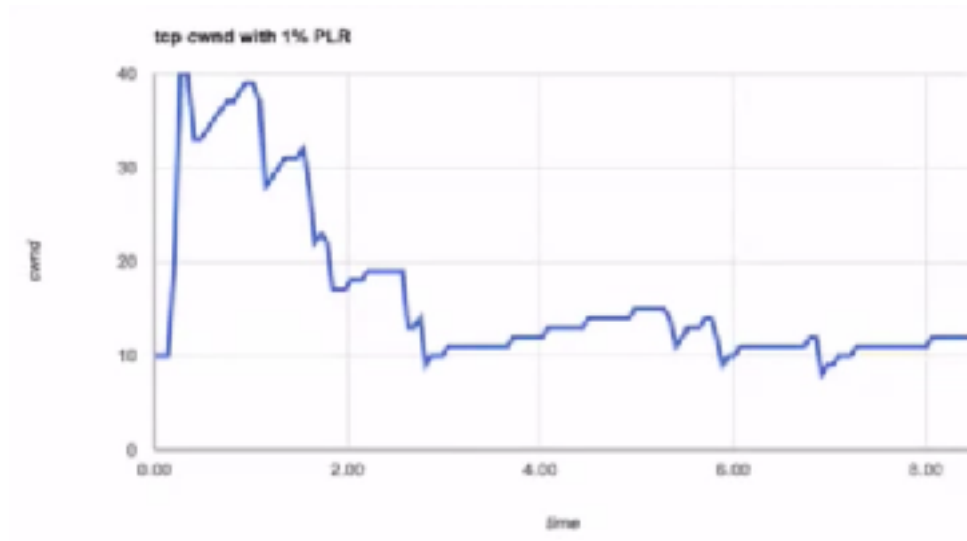
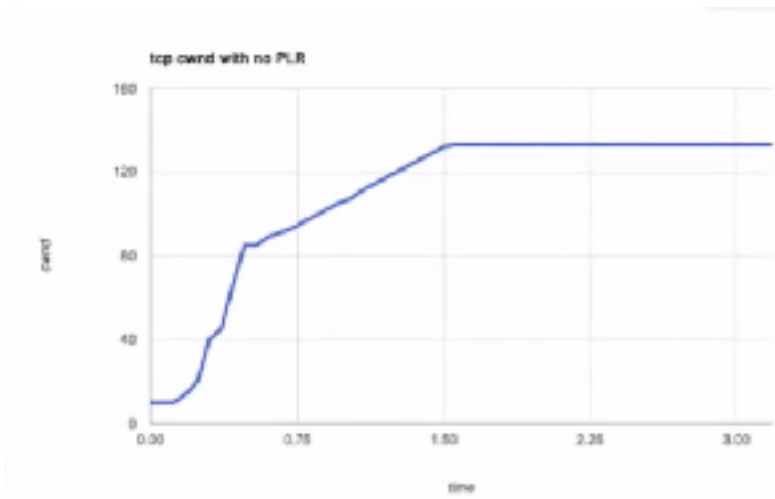
# A single TCP connection

---



# A single TCP connection

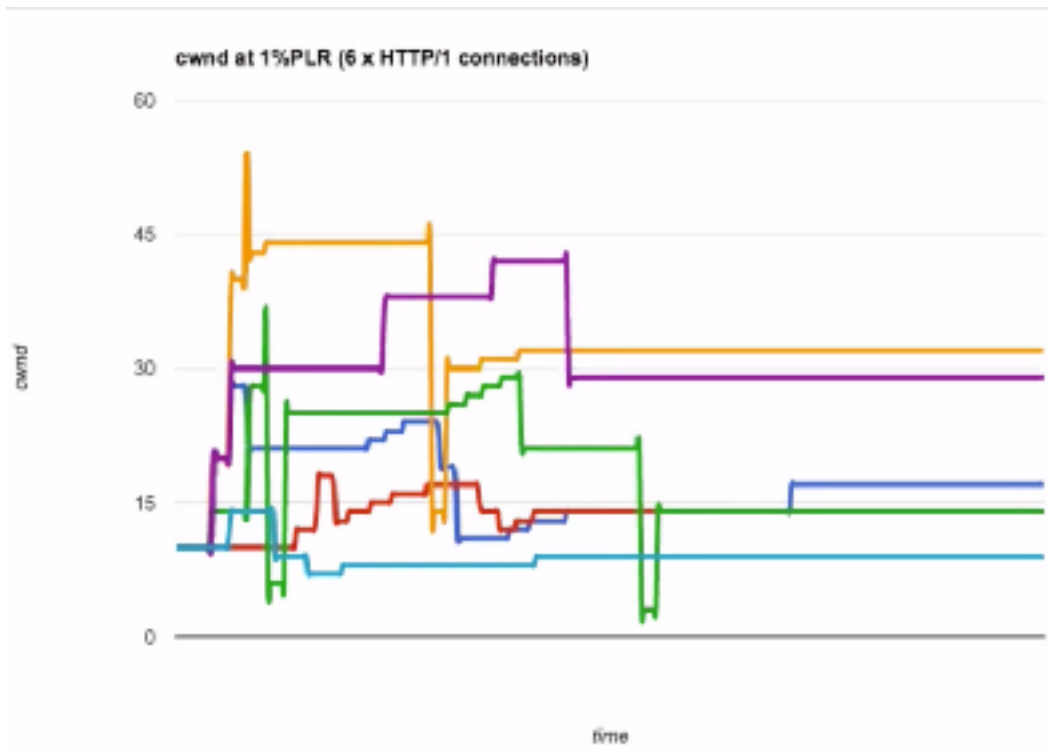
---



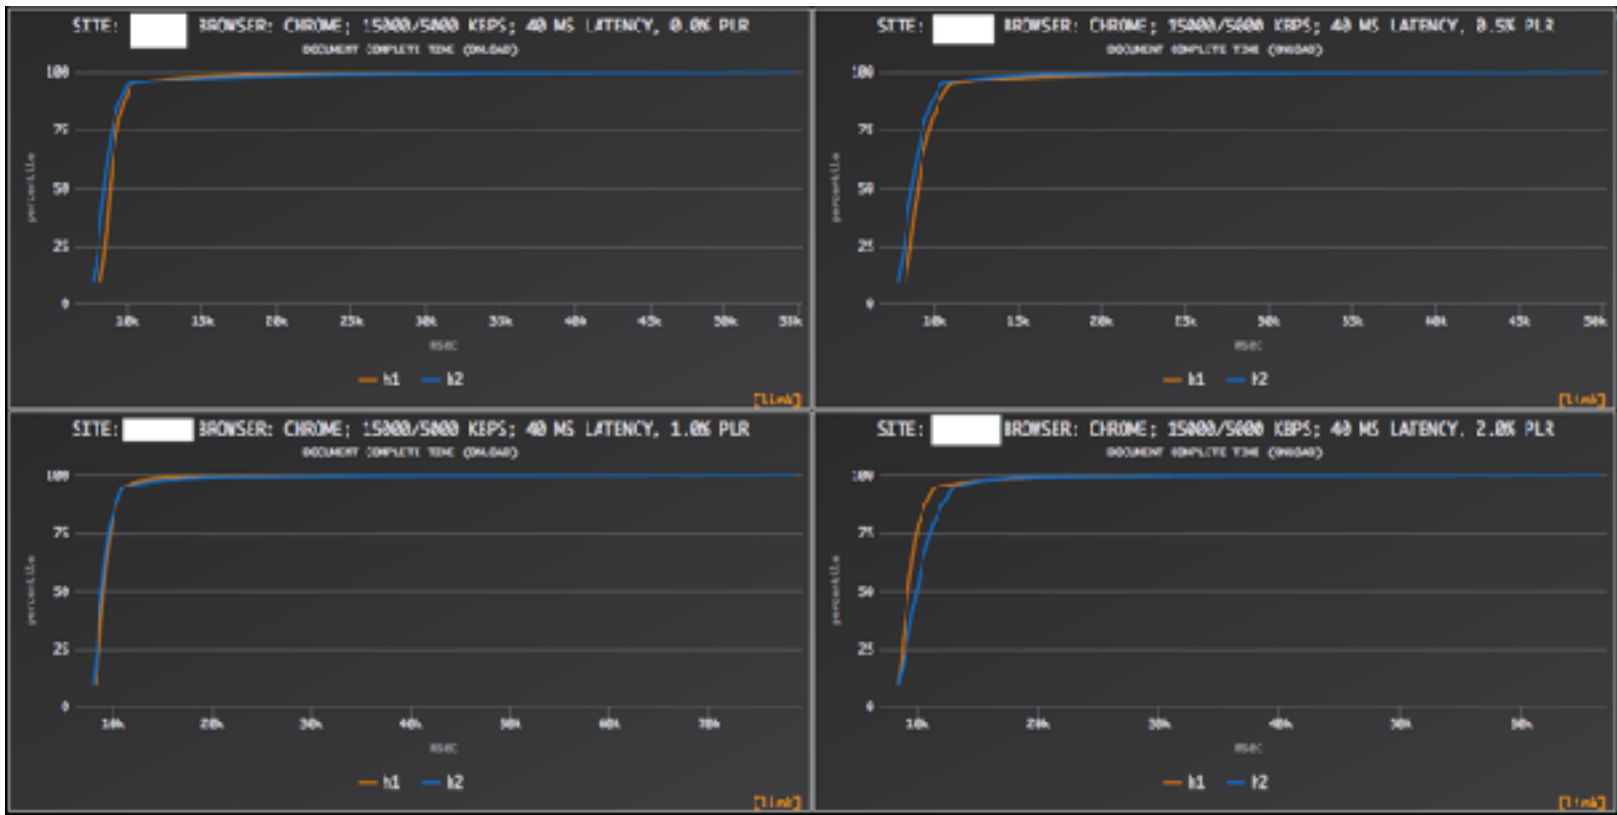


# HTTP/1.1 with packet loss

Could sharding actually help even in HTTP/2



# Real web page

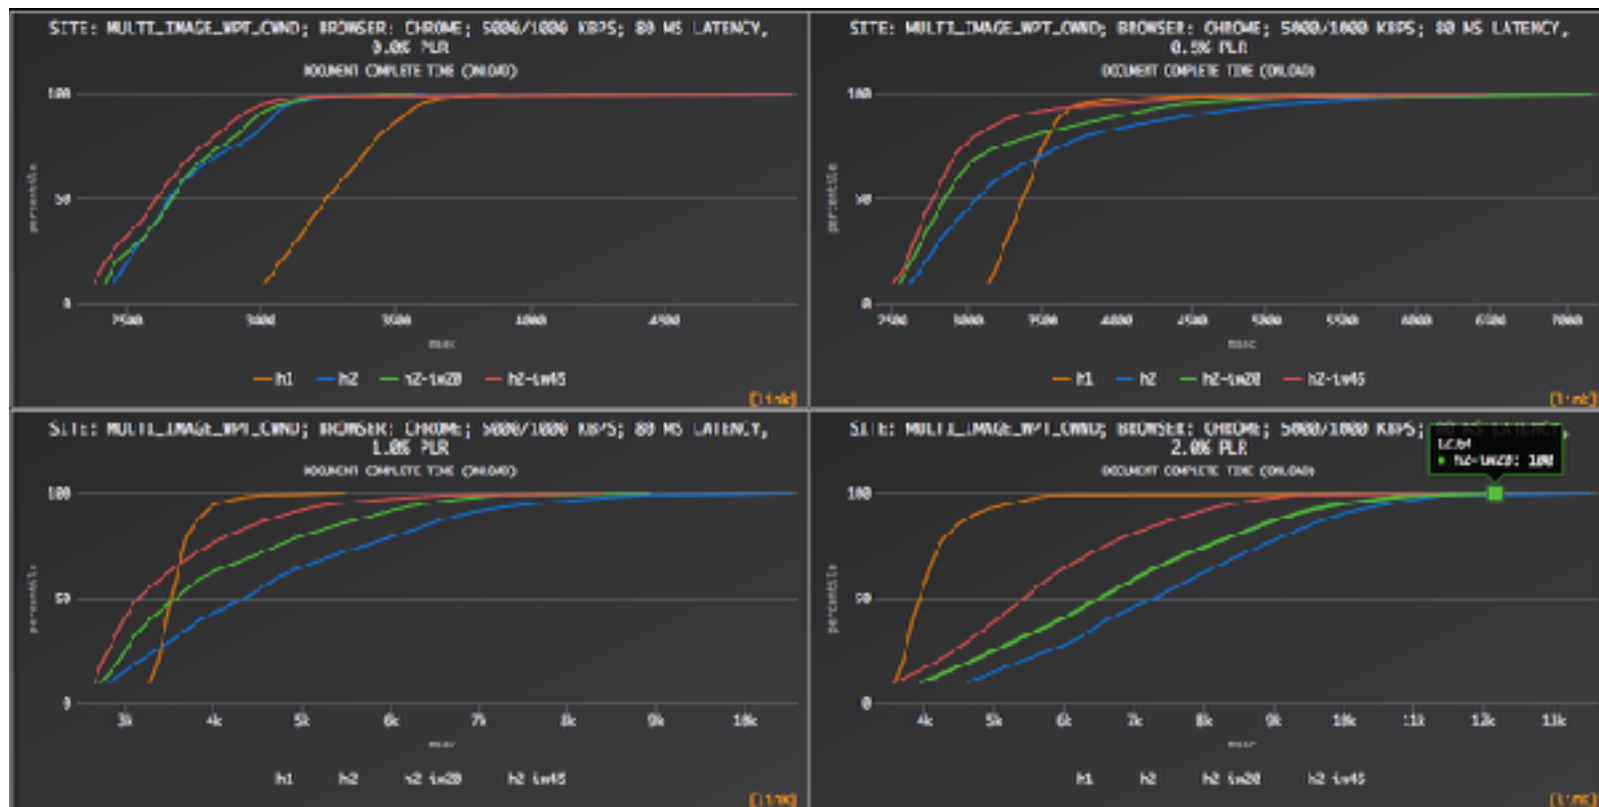


# ICWND experiment

---

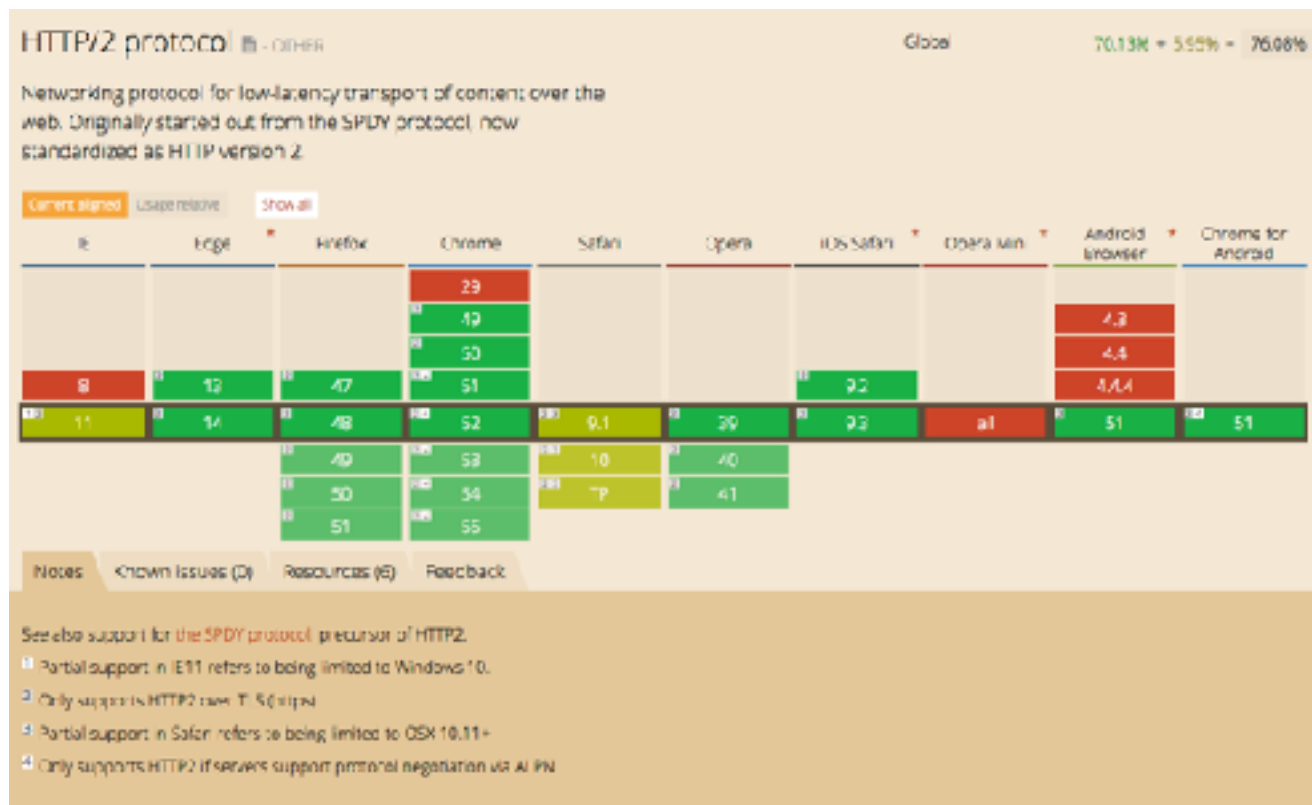
- HTTP/1.1 has 6 times ICWND of 10
- What if we gave the H2 connection a ICWND of 60?

# ICWND experiment



# HTTP/2 implementations

# Browsers



# Servers

## Server software

- Apache 2.4.12 supports HTTP/2 via the module `mod_http2`,<sup>[20]</sup> although appropriate patches must be added to the source code of the server in order for it to support that module. As of Apache 2.4.13 all patches are included in the main Apache source tree, although the module itself was renamed `mod_http2`.<sup>[21]</sup> Old versions of (32-bit) were supported via the module `mod_ssl`,<sup>[22]</sup> however the development of the module has stopped.<sup>[23]</sup>
- Apache Traffic Server supports HTTP/2 with version 3.2 and 3.3 with a configuration change.<sup>[24]</sup>
- Aardite Traffic Server supports HTTP/2.<sup>[25]</sup>
- Caddy supports HTTP/2.<sup>[26]</sup>
- Citrix NetScaler 11.2 supports HTTP/2.<sup>[27]</sup>
- SUSE Linux supports HTTP/2.<sup>[28]</sup>
- F5 BIG-IP Local Traffic Manager 11.6 supports HTTP/2.<sup>[29]</sup>
- IIS was built from the ground up for HTTP/2 support.<sup>[30]</sup>
- Jetty 9.4 supports HTTP/2.<sup>[31]</sup>
- LiteSpeed Web Server 5.0 supports HTTP/2.<sup>[32]</sup>
- Microsoft IIS supports HTTP/2 in Windows 10<sup>[33]</sup> and Windows Server 2016.<sup>[34]</sup>
- Nginx 1.10 supports HTTP/2.<sup>[35]</sup>
- agent 1.1.0 supports HTTP/2.<sup>[36]</sup>
- nginx 1.10 supports HTTP/2.<sup>[37]</sup>
- OpenLiteSpeed 1.3.11 and 1.3.8 supports HTTP/2.<sup>[38]</sup>
- Prerogative 5 supports HTTP/2.<sup>[39]</sup>
- Redware AIXON 4.0 supports HTTP/2.<sup>[40]</sup>
- ShimmyNet was built from the ground up for HTTP/2 support.<sup>[41]</sup>
- Vert.x 3.0 supports HTTP/2.<sup>[42]</sup>
- Warp (Haskell web server, used by reddit in 2015) supports HTTP/2.<sup>[43]</sup>
- WebKit 9 supports HTTP/2.<sup>[44]</sup>

## Content delivery networks

- Akamai is the first major CDN to support HTTP/2 and HTTP/2 Server Push. <http://akamai.com/akamai/akamai-http2-implementations>, including Server Push.
- CloudFlare supports HTTP/2 using nghttp (August 26, 2015). <http://blog.cloudflare.com/http2-early-preview/> is a demonstration of CloudFlare's HTTP/2 implementation.
- CloudFlare supports HTTP/2 using nghttp with CloudFlare as a fallback for browsers without support, whilst maintaining all security and performance services.<sup>[45]</sup> CloudFlare are the first major CDN to support HTTP/2 Server Push.<sup>[46]</sup>
- Fastly<sup>®</sup> supports HTTP/2 including Server Push.<sup>[47]</sup>
- Imperva Incapsula CDN supports HTTP/2.<sup>[48]</sup> <http://incapsula.com/press/press-releases/incapsula-http2-implementation>. The implementation includes support for WAF and DDoS mitigation features as well.
- KeyCDN supports HTTP/2 using nghttp (October 6, 2015). HTTP/2 Test 2 is a test page to verify if your server supports HTTP/2.

# HTTP/2 tooling





# HTTP/2 and SPDY indicator

offered by cauldrip

★★★★★ (162)

[Developer Tools](#)

70,843 users

ADDED TO CHROME



OVERVIEW

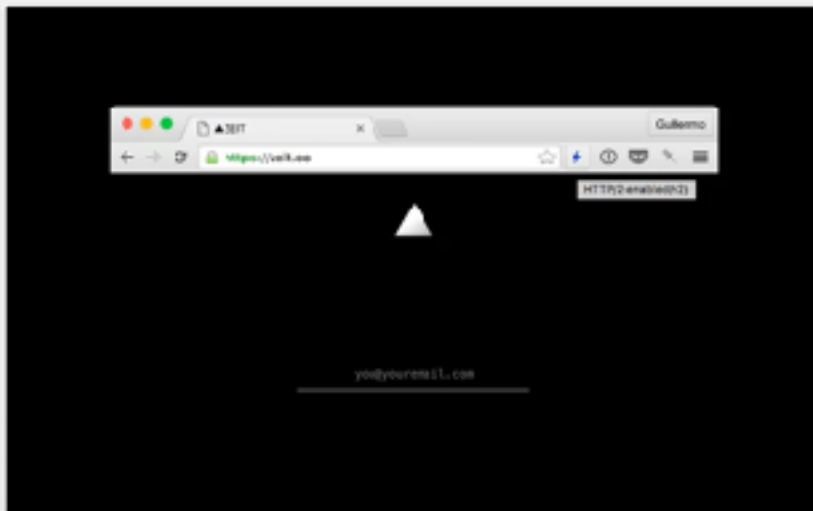
REVIEWS

SUPPORT

RELATED

G-4

444



Compatible with your device

An indicator button for HTTP/2, SPDY and QUIC support by each website.

HTTP/2 (based on Google's SPDY) makes the information exchange between browser and server significantly more performant. Websites and applications that upgrade their infrastructure to support them are at a clear advantage.

This extension shows a lightning icon in the address bar that's blue if the page is HTTP/2 enabled, green if the page is SPDY enabled, and gray if neither is available.

[Privacy Policy](#) [Terms of Use](#) [About](#) [Feedback](#)

[Report Abuse](#)

## Additional information

Version: 1.0.0

Updated: March 11, 2016

Size: 42.79KB

Language: English

USERS OF THIS EXTENSION HAVE ALSO USED



JSONView

★★★★★ (2973)



ColorZilla

★★★★★ (1083)



WharFont

★★★★★ (1000)



Capture Webpage

screenshot entirely

★★★★★ (15147)



# Chrome's net internals

[illegible]

## h2a - a mitmproxy

[illegible]

# curl - since 7.43.0

---

```
$ curl --http2 -I https://www.fastly.com
HTTP/2.0 200
date:Thu, 18 Aug 2016 22:54:47 GMT
content-language:en
content-type:text/html; charset=utf-8
link:<https://www.fastly.com/>; rel="canonical",<https://www.fastly.com/>; rel="shortlink"
via:1.1 varnish
x-content-type-options:nosniff
x-generator:Drupal 7 (http://drupal.org)
last-modified:Thu, 18 Aug 2016 01:51:46 GMT
via:1.1 varnish
```

# nghttp2

```
0 $ nghttp -rv https://www.fastly.com
[ 0.011] Connected
The negotiated protocol: h2
[ 0.045] send SETTINGS frame <length=12, flags=0x00, stream_id=0>
(nlv=2)
[SETTINGS_MAX_CONCURRENT_STREAMS(0x01):100]
[SETTINGS_INITIAL_WINDOW_SIZE(0x04):65535]
[ 0.045] send PRIORITY frame <length=5, flags=0x00, stream_id=5>
(dep_stream_id=0, weight=201, exclusive=0)
[ 0.045] send PRIORITY frame <length=5, flags=0x00, stream_id=5>
(dep_stream_id=0, weight=101, exclusive=0)
[ 0.045] send PRIORITY frame <length=5, flags=0x00, stream_id=7>
(dep_stream_id=0, weight=1, exclusive=0)
[ 0.045] send PRIORITY frame <length=5, flags=0x00, stream_id=9>
(dep_stream_id=7, weight=1, exclusive=0)
[ 0.045] send PRIORITY frame <length=5, flags=0x00, stream_id=11>
(dep_stream_id=9, weight=1, exclusive=0)
[ 0.045] send HEADERS frame <length=39, flags=0x25, stream_id=13>
; END_STREAM | END_HEADERS | PRIORITY
(padlen=0, dep_stream_id=11, weight=16, exclusive=0)
; Open new stream
:method: GET
:path: /
:scheme: https
:authority: www.fastly.com
:accept: */*
:accept-encoding: gzip, deflate
:user-agent: nghttp2/1.13.0
[ 0.050] recv SETTINGS frame <length=12, flags=0x00, stream_id=0>
(nlv=2)
[SETTINGS_MAX_CONCURRENT_STREAMS(0x01):100]
[SETTINGS_INITIAL_WINDOW_SIZE(0x04):16777216]
[ 0.050] recv SETTINGS frame <length=0, flags=0x01, stream_id=0>
; ACK
(nlv=0)
[ 0.050] send SETTINGS frame <length=0, flags=0x01, stream_id=0>
; ACK
(nlv=0)
```

# Language support

---

- Python
  - <https://github.com/python-hyper/hyper-h2>
- Ruby
  - <https://github.com/igrigorik/http-2>
- Golang (part of the standard distribution)
- C (nghttp2, libh2o)

# Future

# QUIC

- Now has a working group, since Oct 2015
  - <https://datatracker.ietf.org/wg/quic/charter/>
- UDP
  - eliminates TCP HOL
- Full control of congestion control algorithms
  - Lesson from slowly evolving TCP stacks



# QUIC

- Chrome
- Server implementations:
  - <https://github.com/google/proto-quic>
  - <https://devsisters.github.io/goquic>
  - <https://github.com/lucas-clemente/quic-go>
- Client
  - <https://github.com/litespeedtech/lisquic-client>
  - <https://github.com/GoogleChromeLabs/cronet-sample>

# To recap

---

- HTTP/2 is really a browser transport protocol
- We've moved blocking to TCP
- We need more real world data to understand HTTP/2 performance



---

# Thank you